

BLOCKCHAIN FUNDAMENTALS

Kostas Christidis • Duke MMCi 2017.06.23

About me

- Kostas Christidis
- Maintainer, Hyperledger Fabric
- Software Engineer, IBM Blockchain
- PhD Candidate, ECE NCSU

Plan of attack

1. Relevant **problem** in the **healthcare** industry
2. Let's try to **solve** it (**forget** about blockchains!)

Relevant problem¹

- Balkanization of health records
- Lack of immediate access to health information
- 20% of preventable medical errors

¹ Source: L. J. Kish and E. J. Topol, "Unpatients—why patients should own their medical data," Nat. Biotechnol., vol. 33, no. 9, pp. 921–924, Sep. 2015.

Why is it that way?

1. No **incentives** for sharing
2. "We'll gladly let you **access** it"
3. No **interoperability** between EHR systems
4. **Easiest** option (keep the data locked)

What should we aim for?

Make patient data easier to access

What's the easiest way to get there?

A **centralized** service

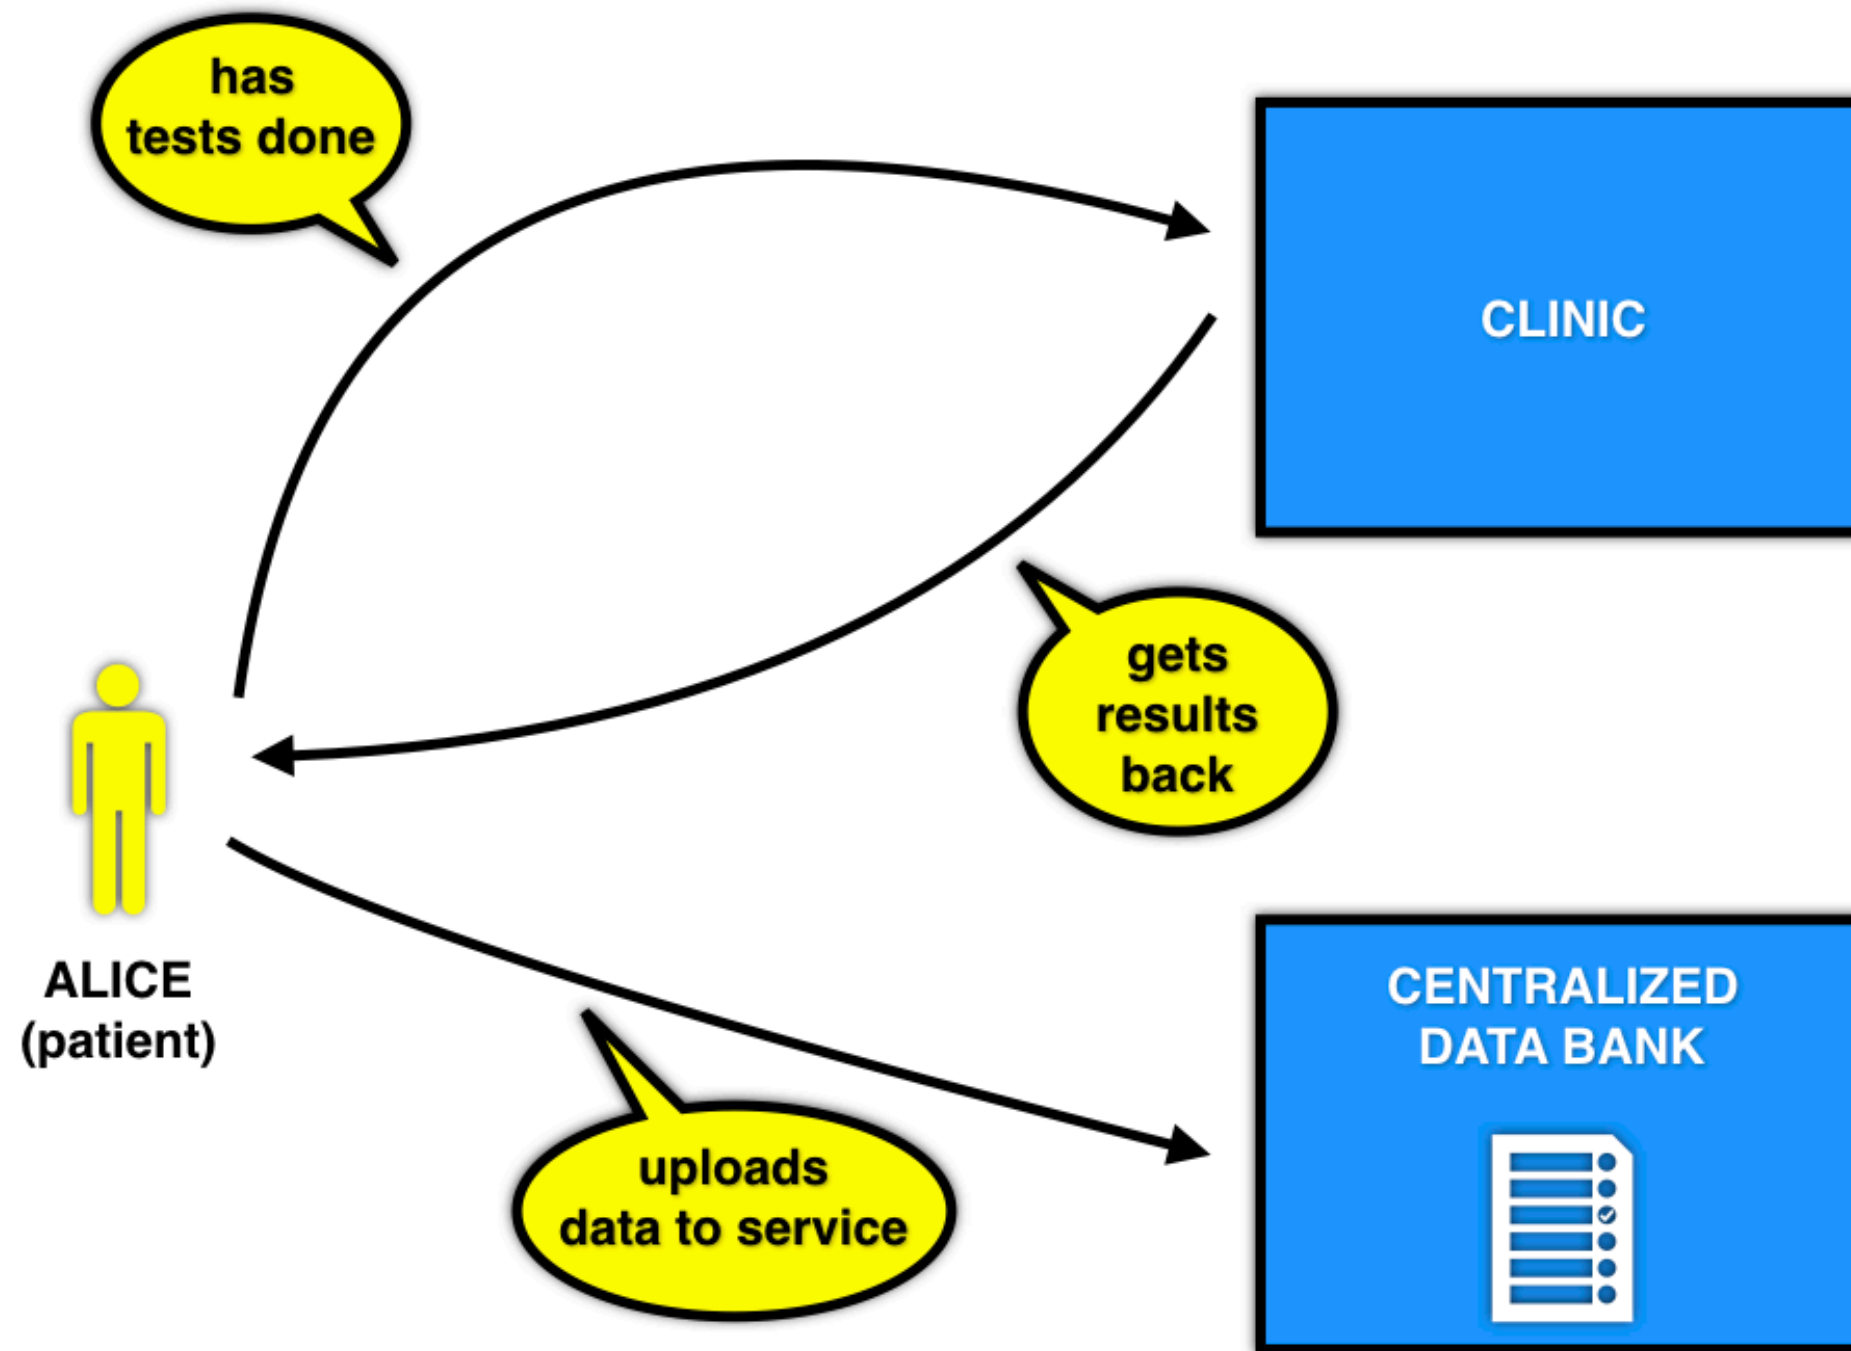
Holds every patient record out there

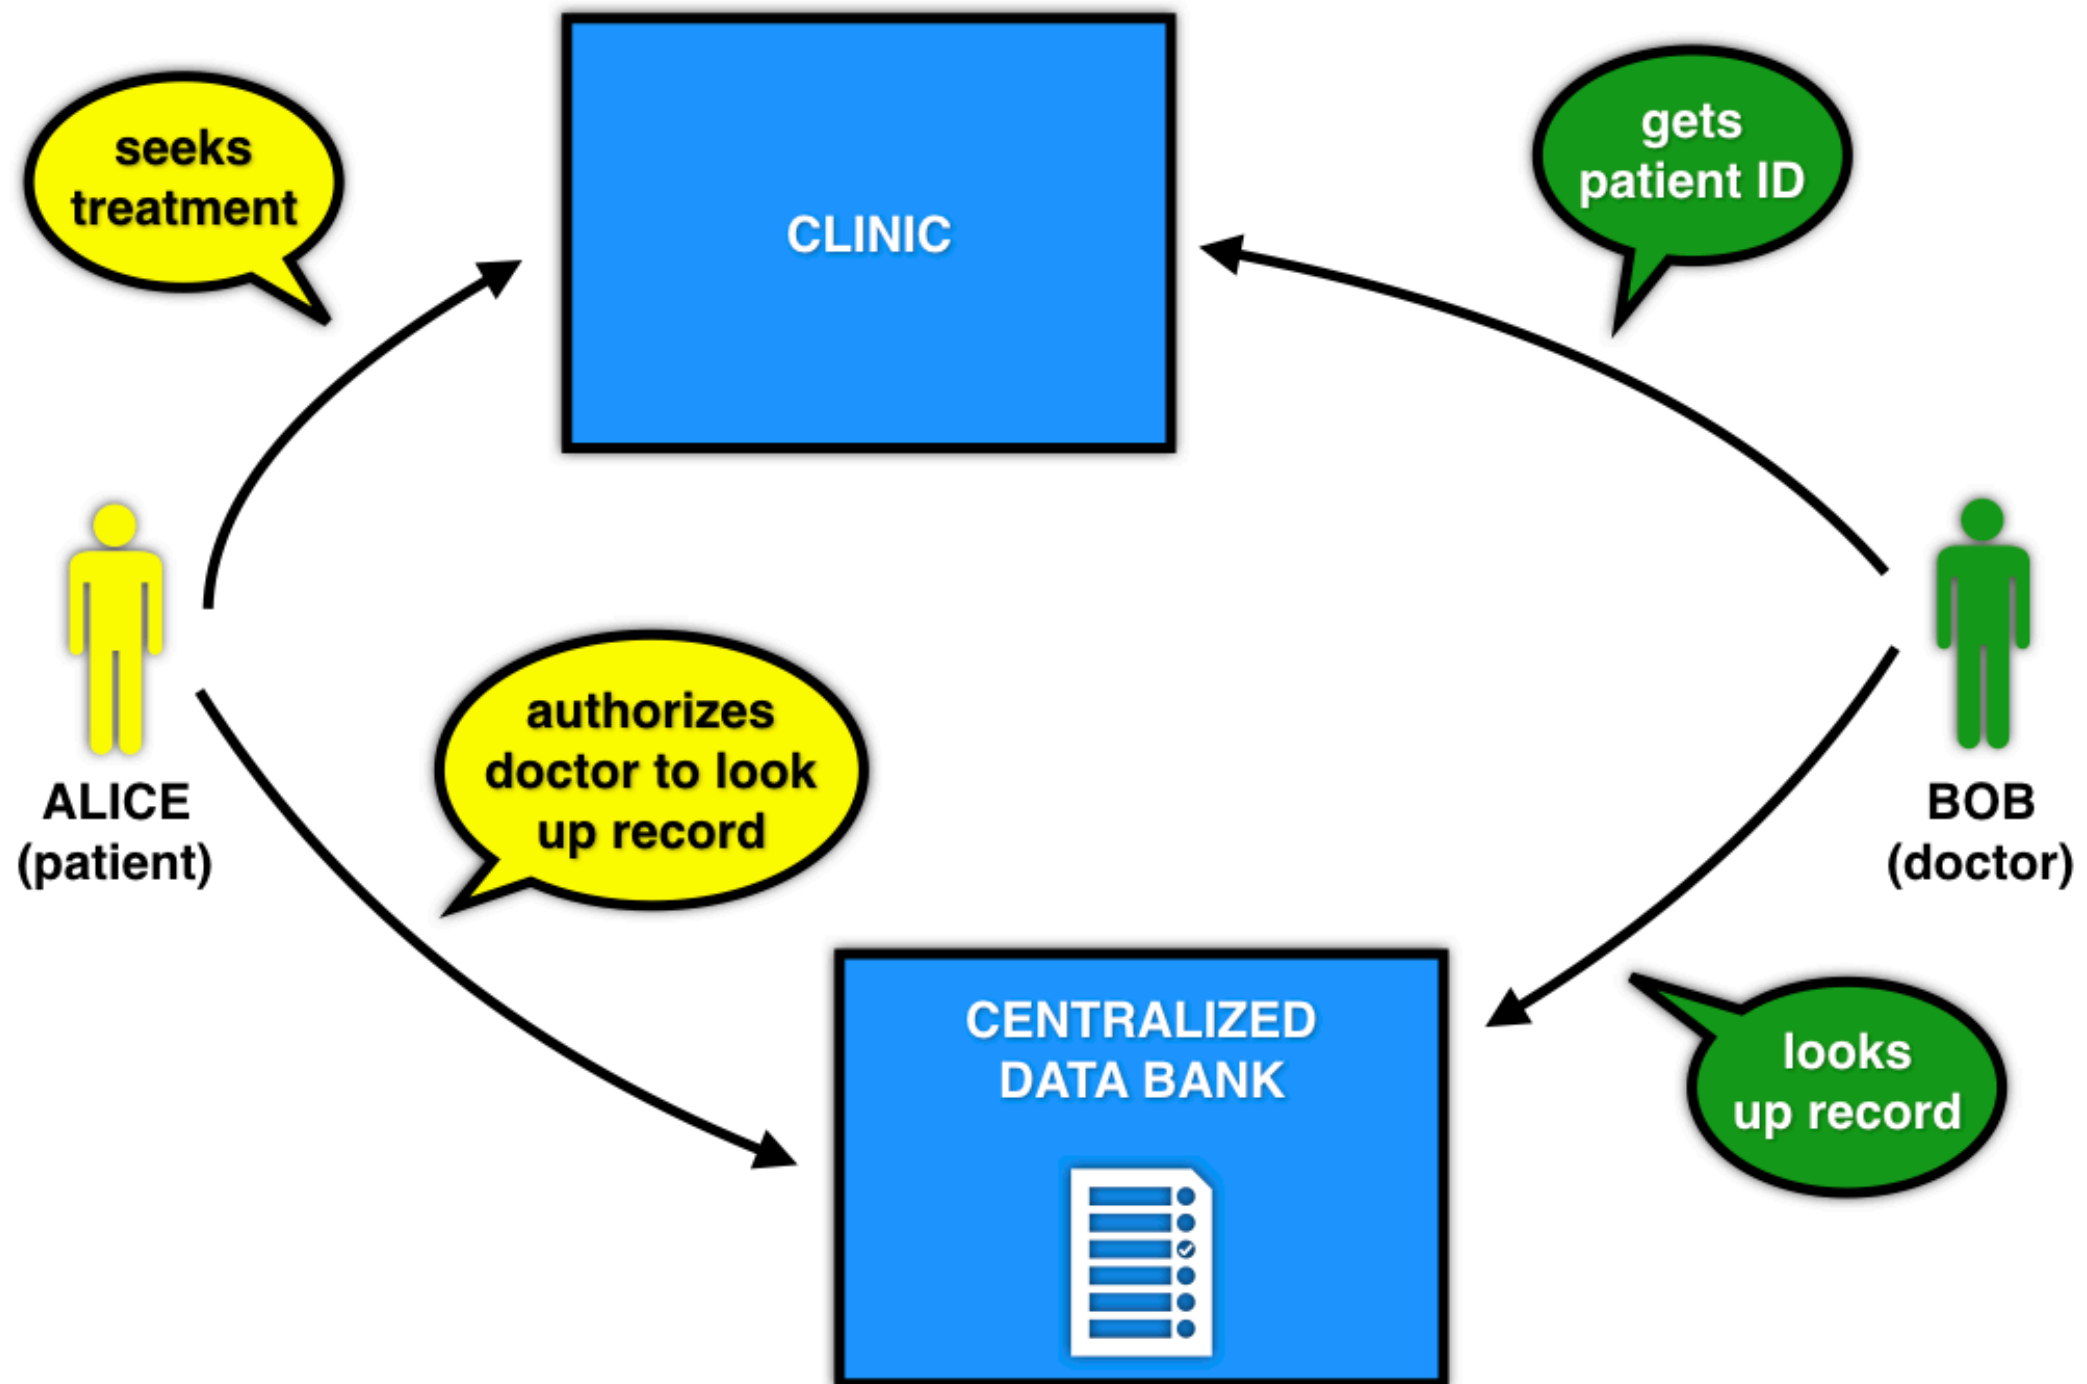
What's the easiest way to get there?

A **centralized** service

Holds every patient record out there

It would look somewhat like this...





This implies a table that looks like this

Patient Name	Record	Who Can Read	Who Can Write
Alice	Blood work @ Duke Hospital (2016/05/01)	Bob	Nobody

This implies a table that looks like this

Patient Name	Record	Who Can Read	Who Can Write
Alice	Blood work @ Duke Hospital (2016/05/01)	Bob	Nobody
John	Shoulder X-rays @ Rex Hospital (2016/02/04)	Carol	Carol

Two concerns with that approach

1. A nightmare if it gets hacked
2. Who gets to run this service?

We need a different approach

If a centralized service won't do...


- Maybe a decentralized solution will?
- Every stakeholder is involved

If a centralized service won't do...

- Maybe a decentralized solution will?
- Every stakeholder is involved

But how would this work exactly?

How would the decentralized service work?



How would the decentralized service work?

1. Every stakeholder* has **a copy of that table**

* Hospital, patient, insurance agent, you-name-it

How would the decentralized service work?

1. Every stakeholder* has **a copy of that table**

* Hospital, patient, insurance agent, you-name-it

Set aside any privacy concerns for now.

In fact, we don't need the table

Transaction logs record all the changes made to the database. [...] appends are the only way to change the log. From this perspective, the contents of the database hold a caching of the latest record values in the logs.

– Pat Helland, "Immutability Changes Everything"

Let's look at an example

Transaction log

$t=0$: $k1 \Rightarrow v1$

Database (table) state

Key

Value

k1

v1

Transaction log

t=0: k1=>v1 | t=1: k2=>v1

Database state

Key	Value
k1	v1
k2	v2

Transaction log

t=0: k1=>v1 | t=1: k2=>v2 | t=2: k1=>v1'

Database state

Key	Value
k1	v1'
k2	v2

The log is all that matters

1. Read it from start to finish
2. You have the most recent state of the table

Again: how would the decentralized service work?

1. Every stakeholder has a copy of that ~~table~~ transaction log

Again: how would the decentralized service work?

1. Every stakeholder has a copy of that ~~table~~ transaction log

How do we ensure the log is kept in sync across participants?

When there's no global ordering, we get:

- Participant A's log: **foo | bar | baz**
- Participant B's log: **foo | baz | bar**

When there's no global ordering, we get:

- Participant A's log: **foo | bar | baz**
- Participant B's log: **foo | baz | bar**

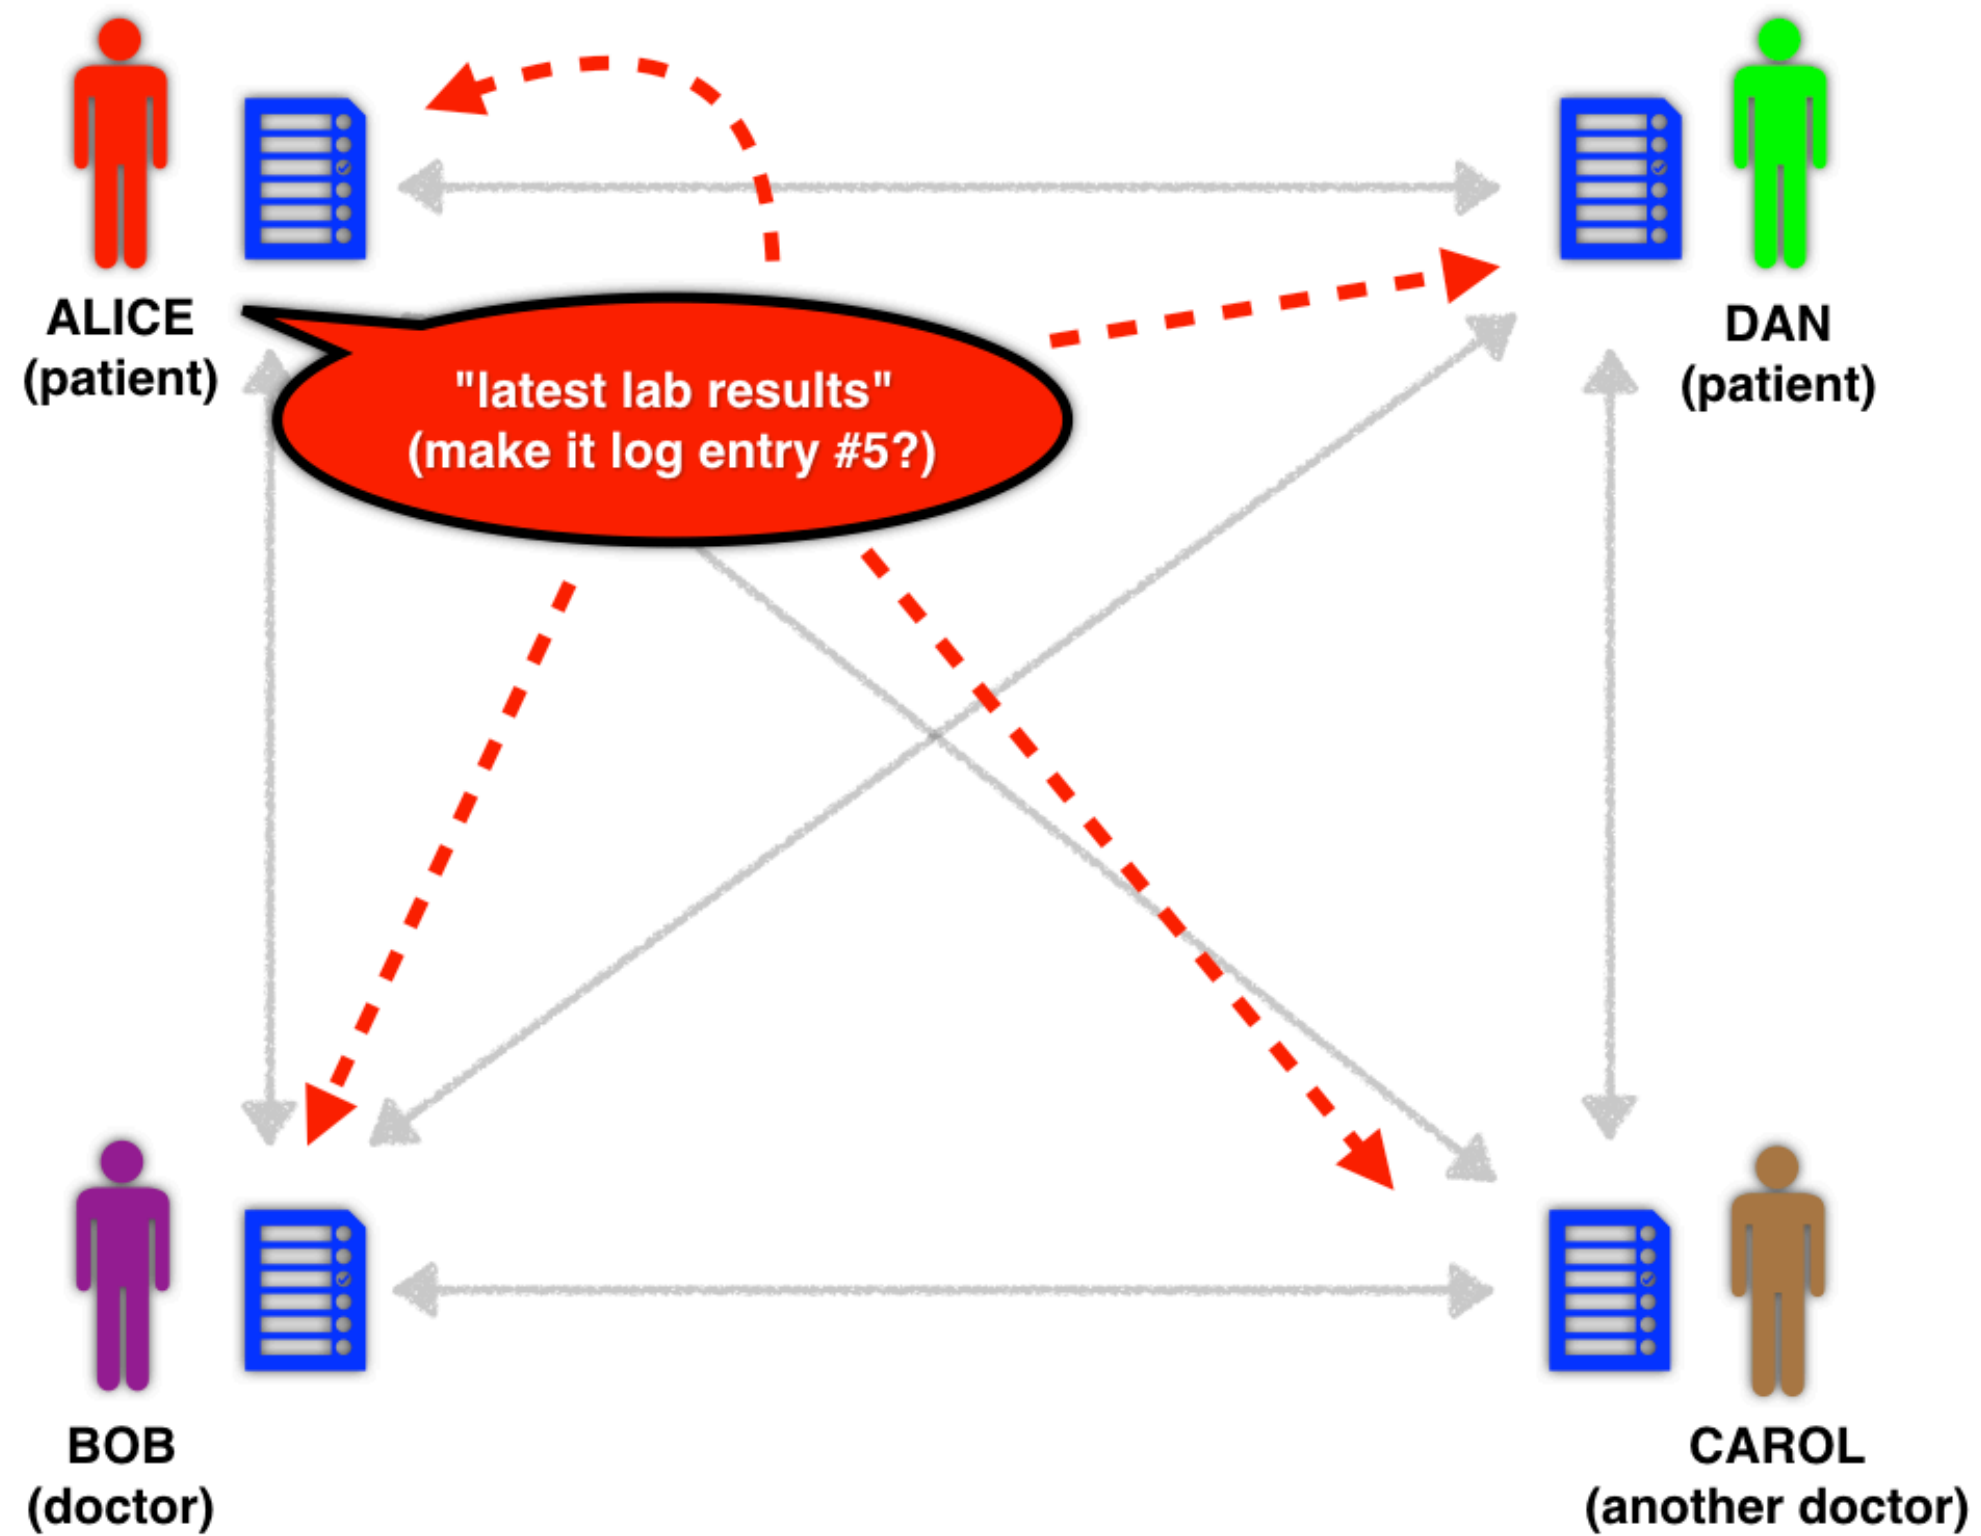
(^^ This is what we call a fork)

How do we order the writes?

- Need a marker
- Before you append to the log, determine its **height**
- Include this info in your broadcasted message

For example, Alice broadcasts this to the network:

*These are my latest lab results. Make it **log entry #5**?*



In fact, we can provide a better marker

A hash function is any function that can be used to map data of arbitrary size to data of fixed size.

Google "sha-256 hash generator"

Input

Output

[foo,bar]

2c26b46b68ffc68ff99b453c1d30
413413422d706483bfa0f98a5e
886266e7ae

[foo,bar,baz]

b94d27b9934d3e08a52e52d7d
a7dabfac484efe37a5380ee9088
f7ace2efcde9

How do we order the writes?

Agree on a hashing function: `calc_hash()`

Assume the log looks like this: `[foo]`

How do we order the writes?

Agree on a hashing function: `calc_hash()`

Assume the log looks like this: `[foo]`

```
hash := calc_hash(latest_log_entry)
```

```
message := <hash, original_message>
```

```
broadcast(message)
```

How do we order the writes? (cont.)

- Log before: `[foo]`
- Log after: `[foo, <calc_hash(foo), original_message>]`

How do we order the writes? (cont.)

- Log before: `[foo]`
- Log after: `[foo, <calc_hash(foo), original_message>]`

Let `bar := <calc_hash(foo), original_message>`

How do we order the writes? (cont.)

- Log before: `[foo]`
- Log after: `[foo, <calc_hash(foo), original_message>]`

Let `bar := <calc_hash(foo), original_message>`

- New log is basically: `[foo, bar]`

How do we order the writes now?

Used to be:

*These are my latest lab results. Make it **log entry #5**?*

How do we order the writes now?

Used to be:

*These are my latest lab results. Make it **log entry #5**?*

Now:

These are my latest lab results. The hash of the log tip is

**2c26b46b68ffc68ff99b453c1d30413413422d706483bfa
0f98a5e886266e7ae.**

Why is this a better marker than log height?

- Acts as a built-in consistency check
- You're not saying: "add to page 50"
- You're saying: "add to page 50 of a book that has this written in the previous 49 pages"

Why is this a better marker than log height?

- Acts as a built-in consistency check
- You're not saying: "add to page 50"
- You're saying: "add to page 50 of a book that has this written in the previous 49 pages"

As we'll see below, hashing combined with digital signatures

Status check

1. Every stakeholder has a copy of the transaction log
2. Writes ordered by including the hash of the tip of the log


Status check


1. Every stakeholder has a copy of the transaction log
2. Writes ordered by including the hash of the tip of the log

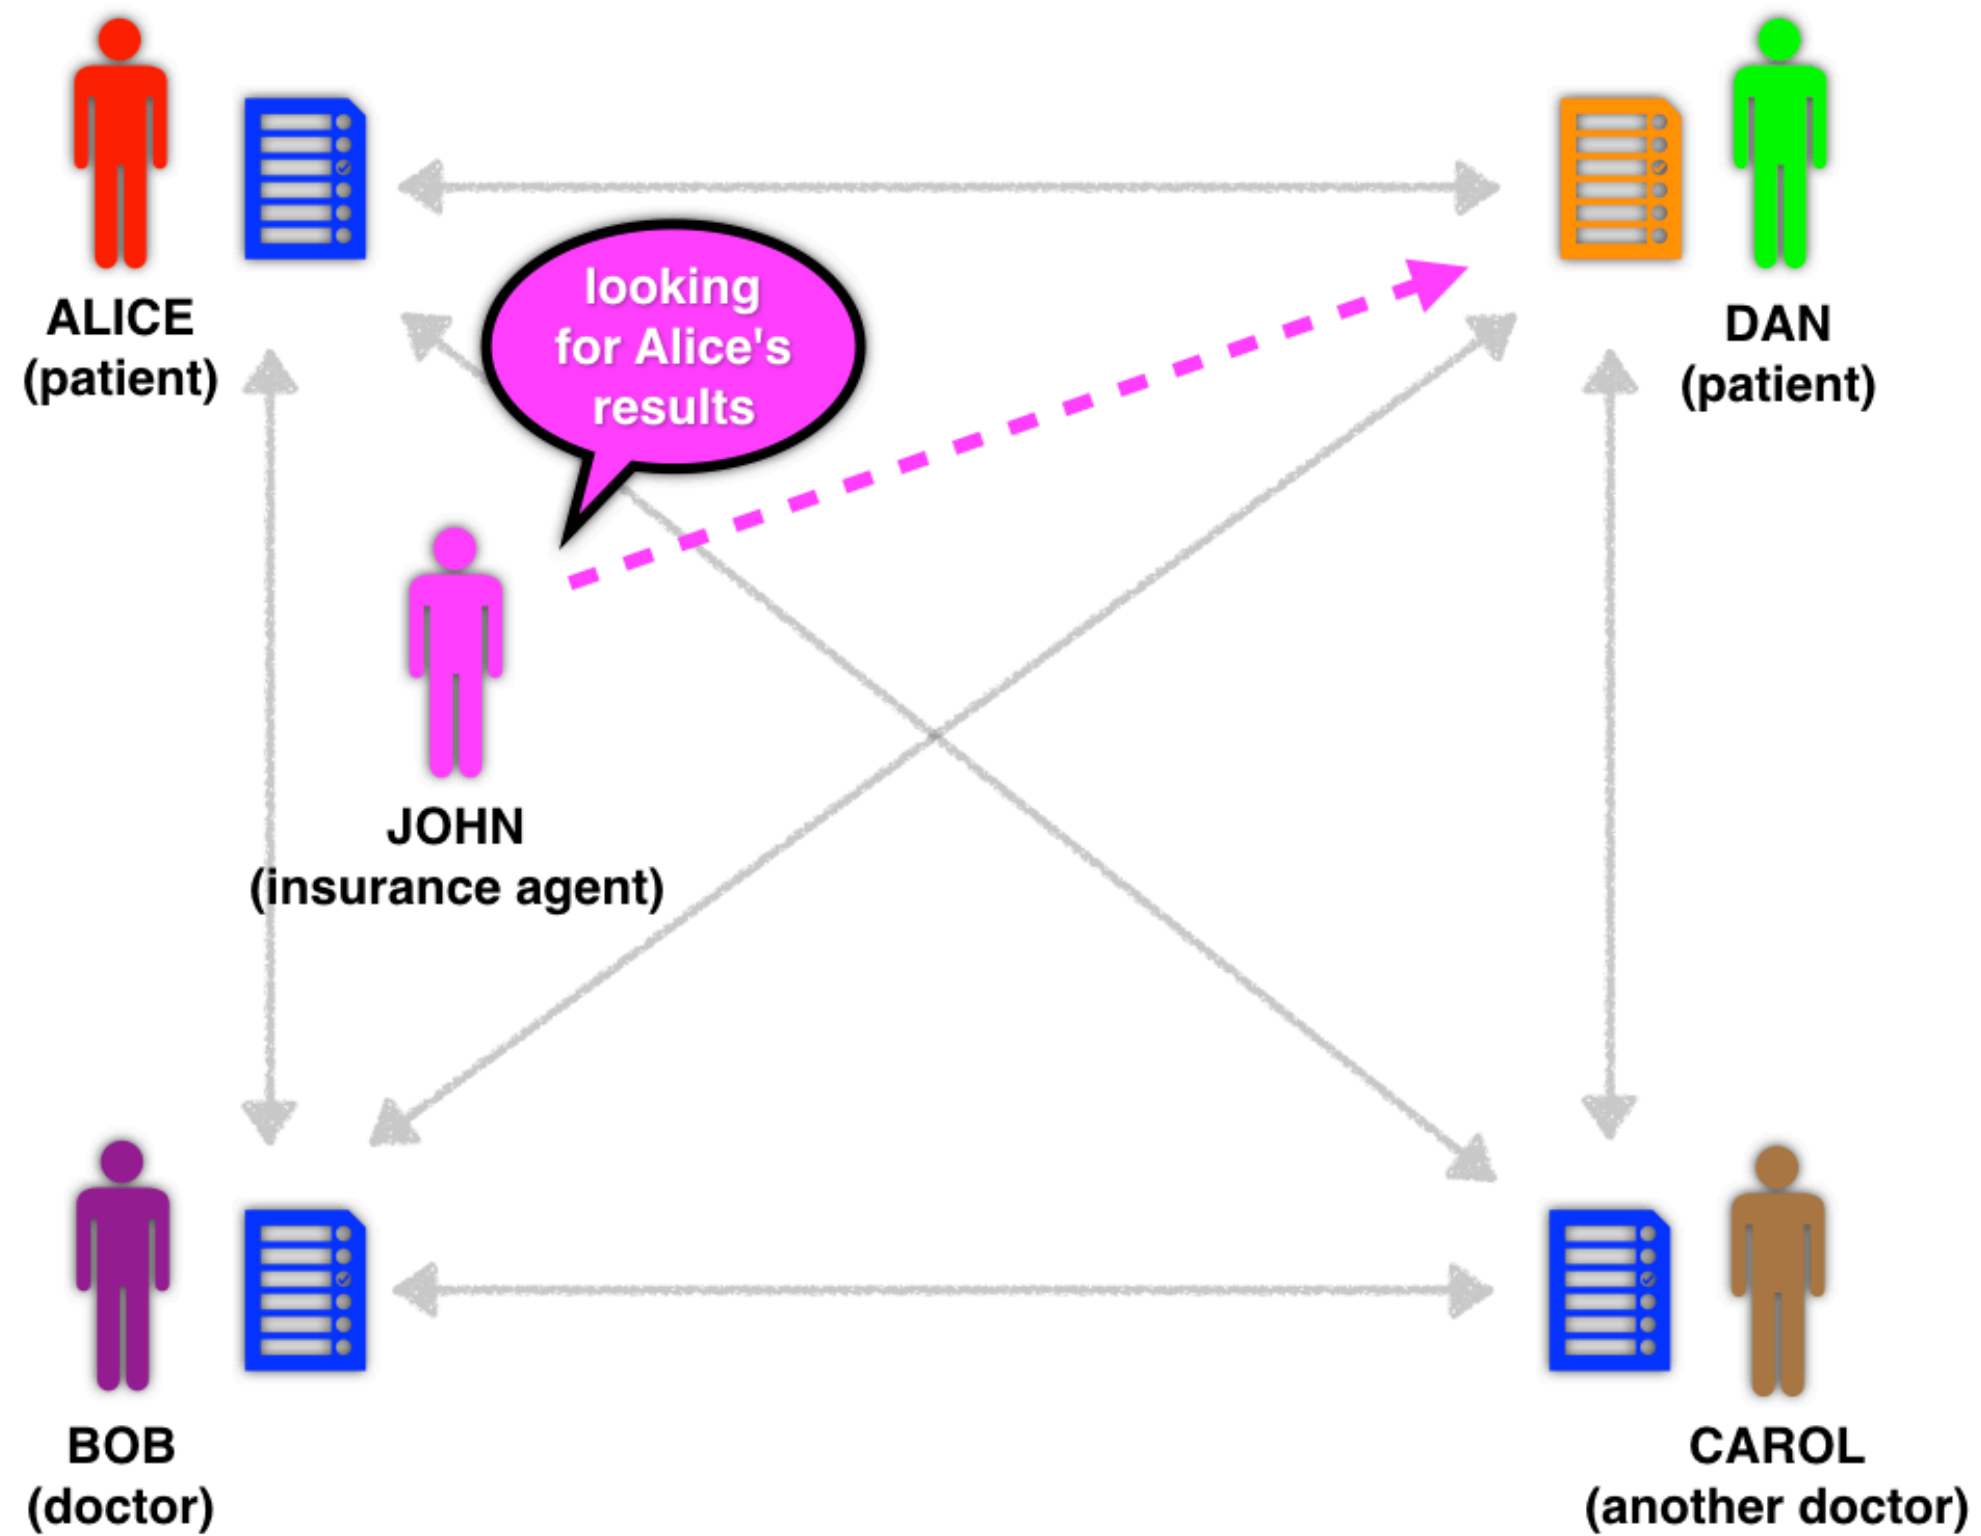
*Ignore collisions for now.
Will address them later.*

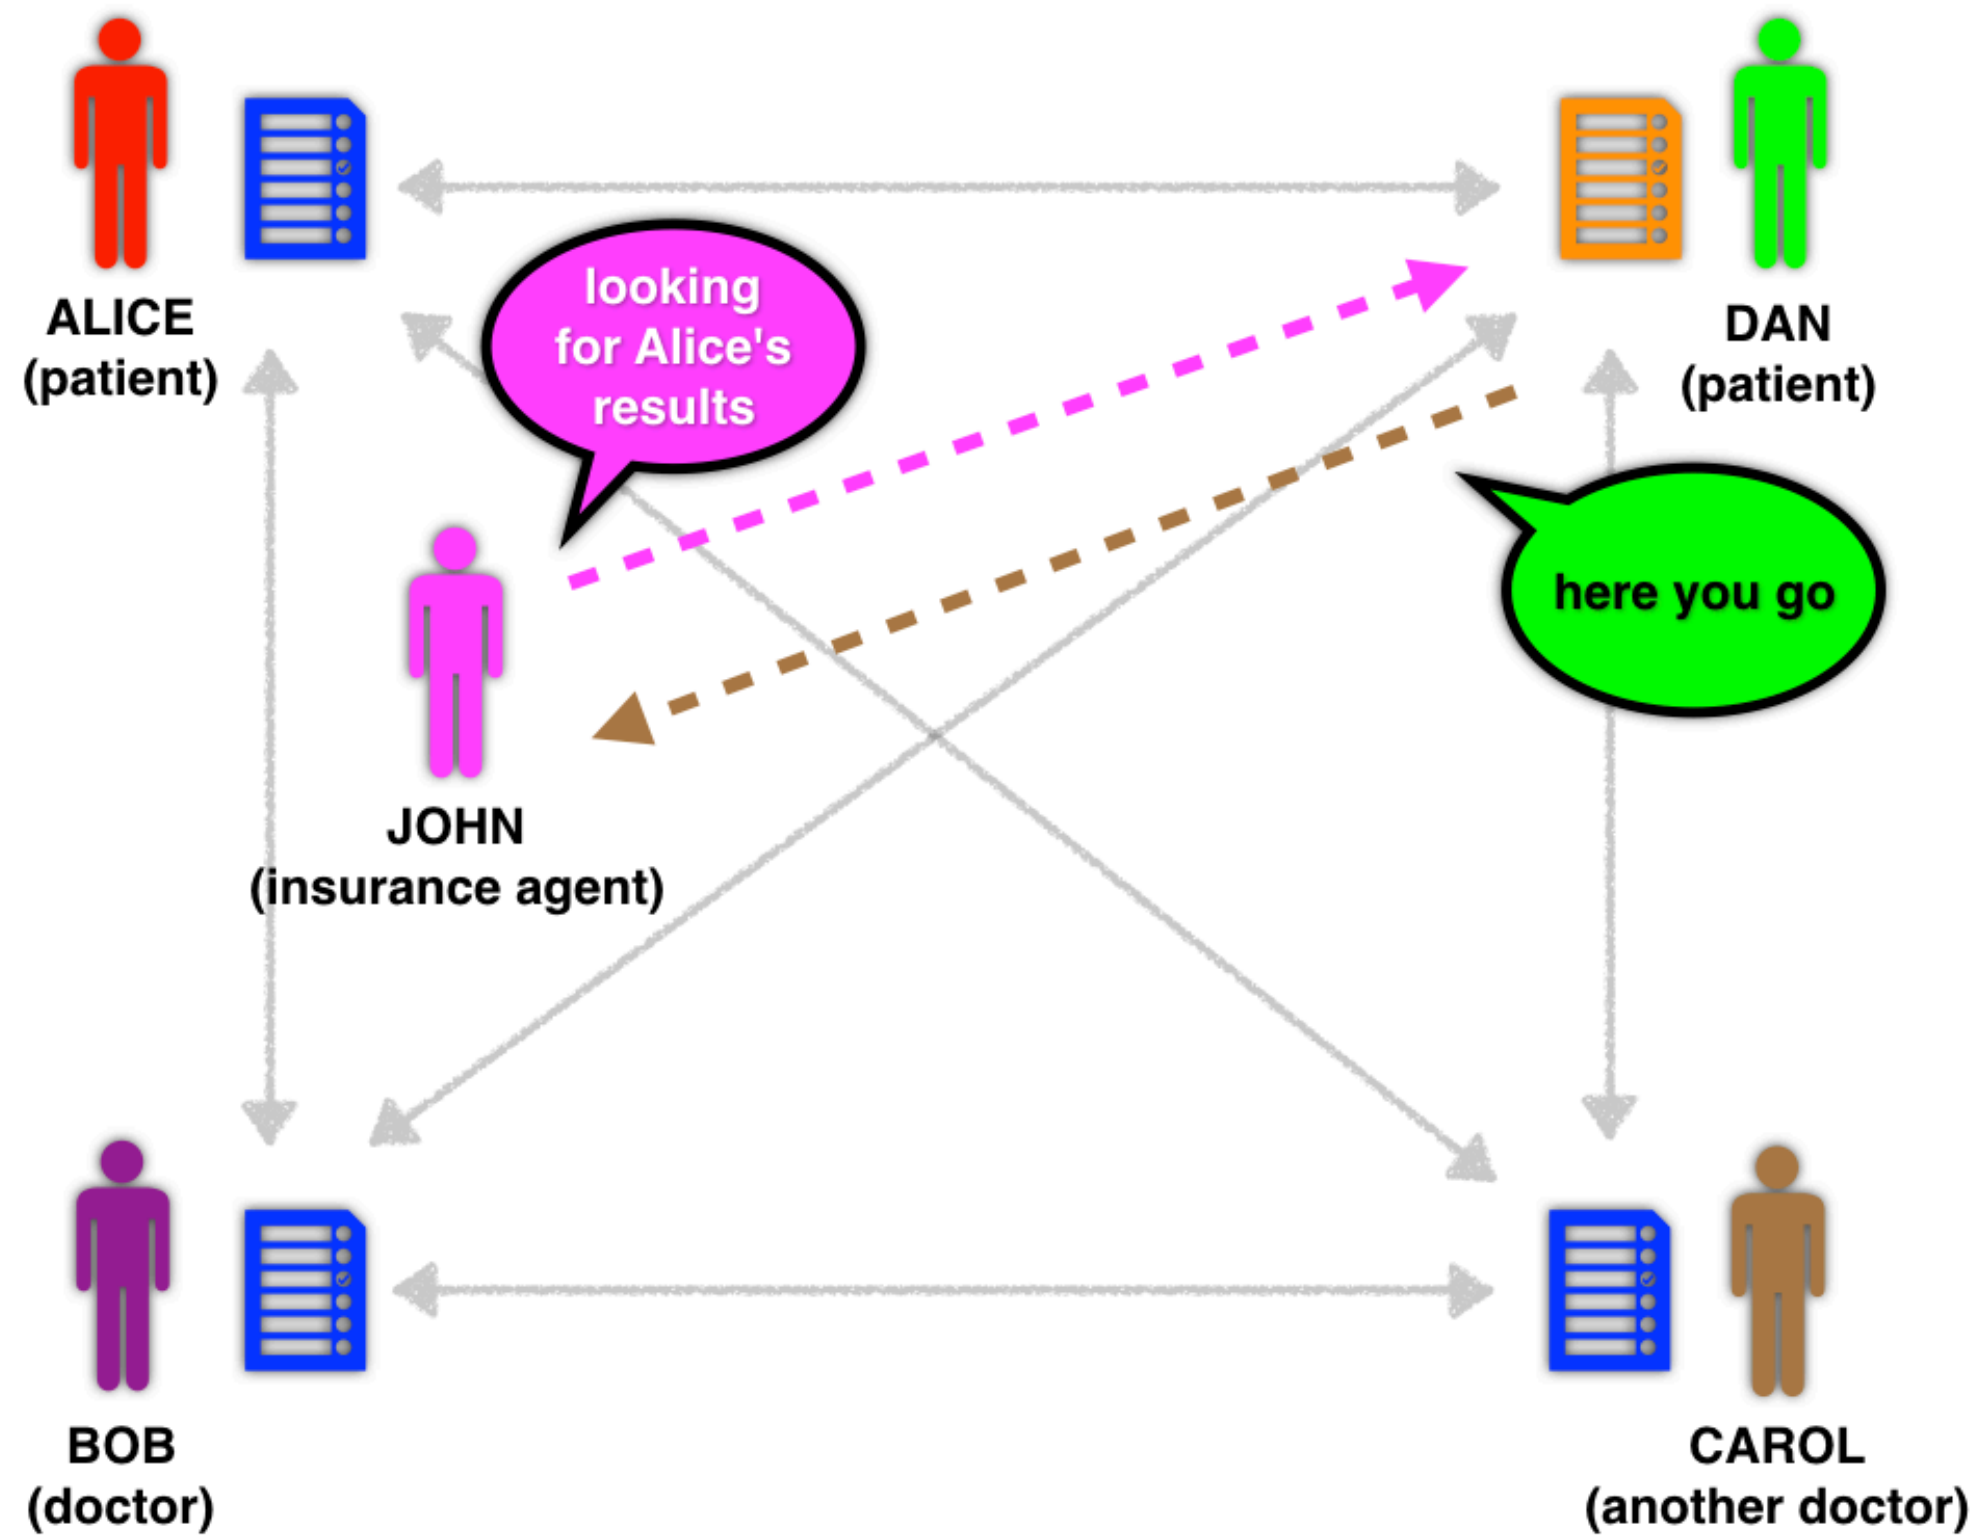
Are we done?

Consider this scenario









Scenario summary

- Alice uploads the results of her latest checkup
- (She's in perfect health)
- Her insurance agent checks them to adjust her premiums
- Reads them from Bob's log
- Bob tampers Alice's update on purpose
- Alice's premiums get adjusted incorrectly

How do we avoid this?

Public key cryptography to the rescue

- Alice has a **known** public key
- Signs her updates with her private key
- Appends signature to her broadcasted message
- Reader looks for that signature

Pseudocode

```
sign := calc_sign(private_key, message)
signed_update = [sign, message]
// ["835cda89", <2c26b46b, "LDL normal">]
```

We now have a table that looks like this

Patient Name	Public Key	Record	Digital Signature	Who Can Read	Who Can Write
Alice	b83da6c0 9893ad	Blood work @ Duke Hospital (2016/05/ 01)	835cda89 7d74c3	Carol	Nobody

We now have a table that looks like this

Patient Name	Public Key	Record	Digital Signature	Who Can Read	Who Can Write
Alice	b83da6c09893ad	Blood work @ Duke Hospital (2016/05/01)	835cda897d74c3	Bob, Carol	Nobody
Alice	b83da6c09893ad	Shoulder X-rays @ Rex Hospital (2016/02/04)	dac089r967856d	Nobody	Nobody

Digital signatures give us:

- Authentication
- Non-repudiation

Status check

1. Every stakeholder has a copy of the transaction log
2. All writes include the hash of the tip of the log for ordering
3. All writes are signed

Haven't addressed yet:

1. Privacy concerns
2. Collisions

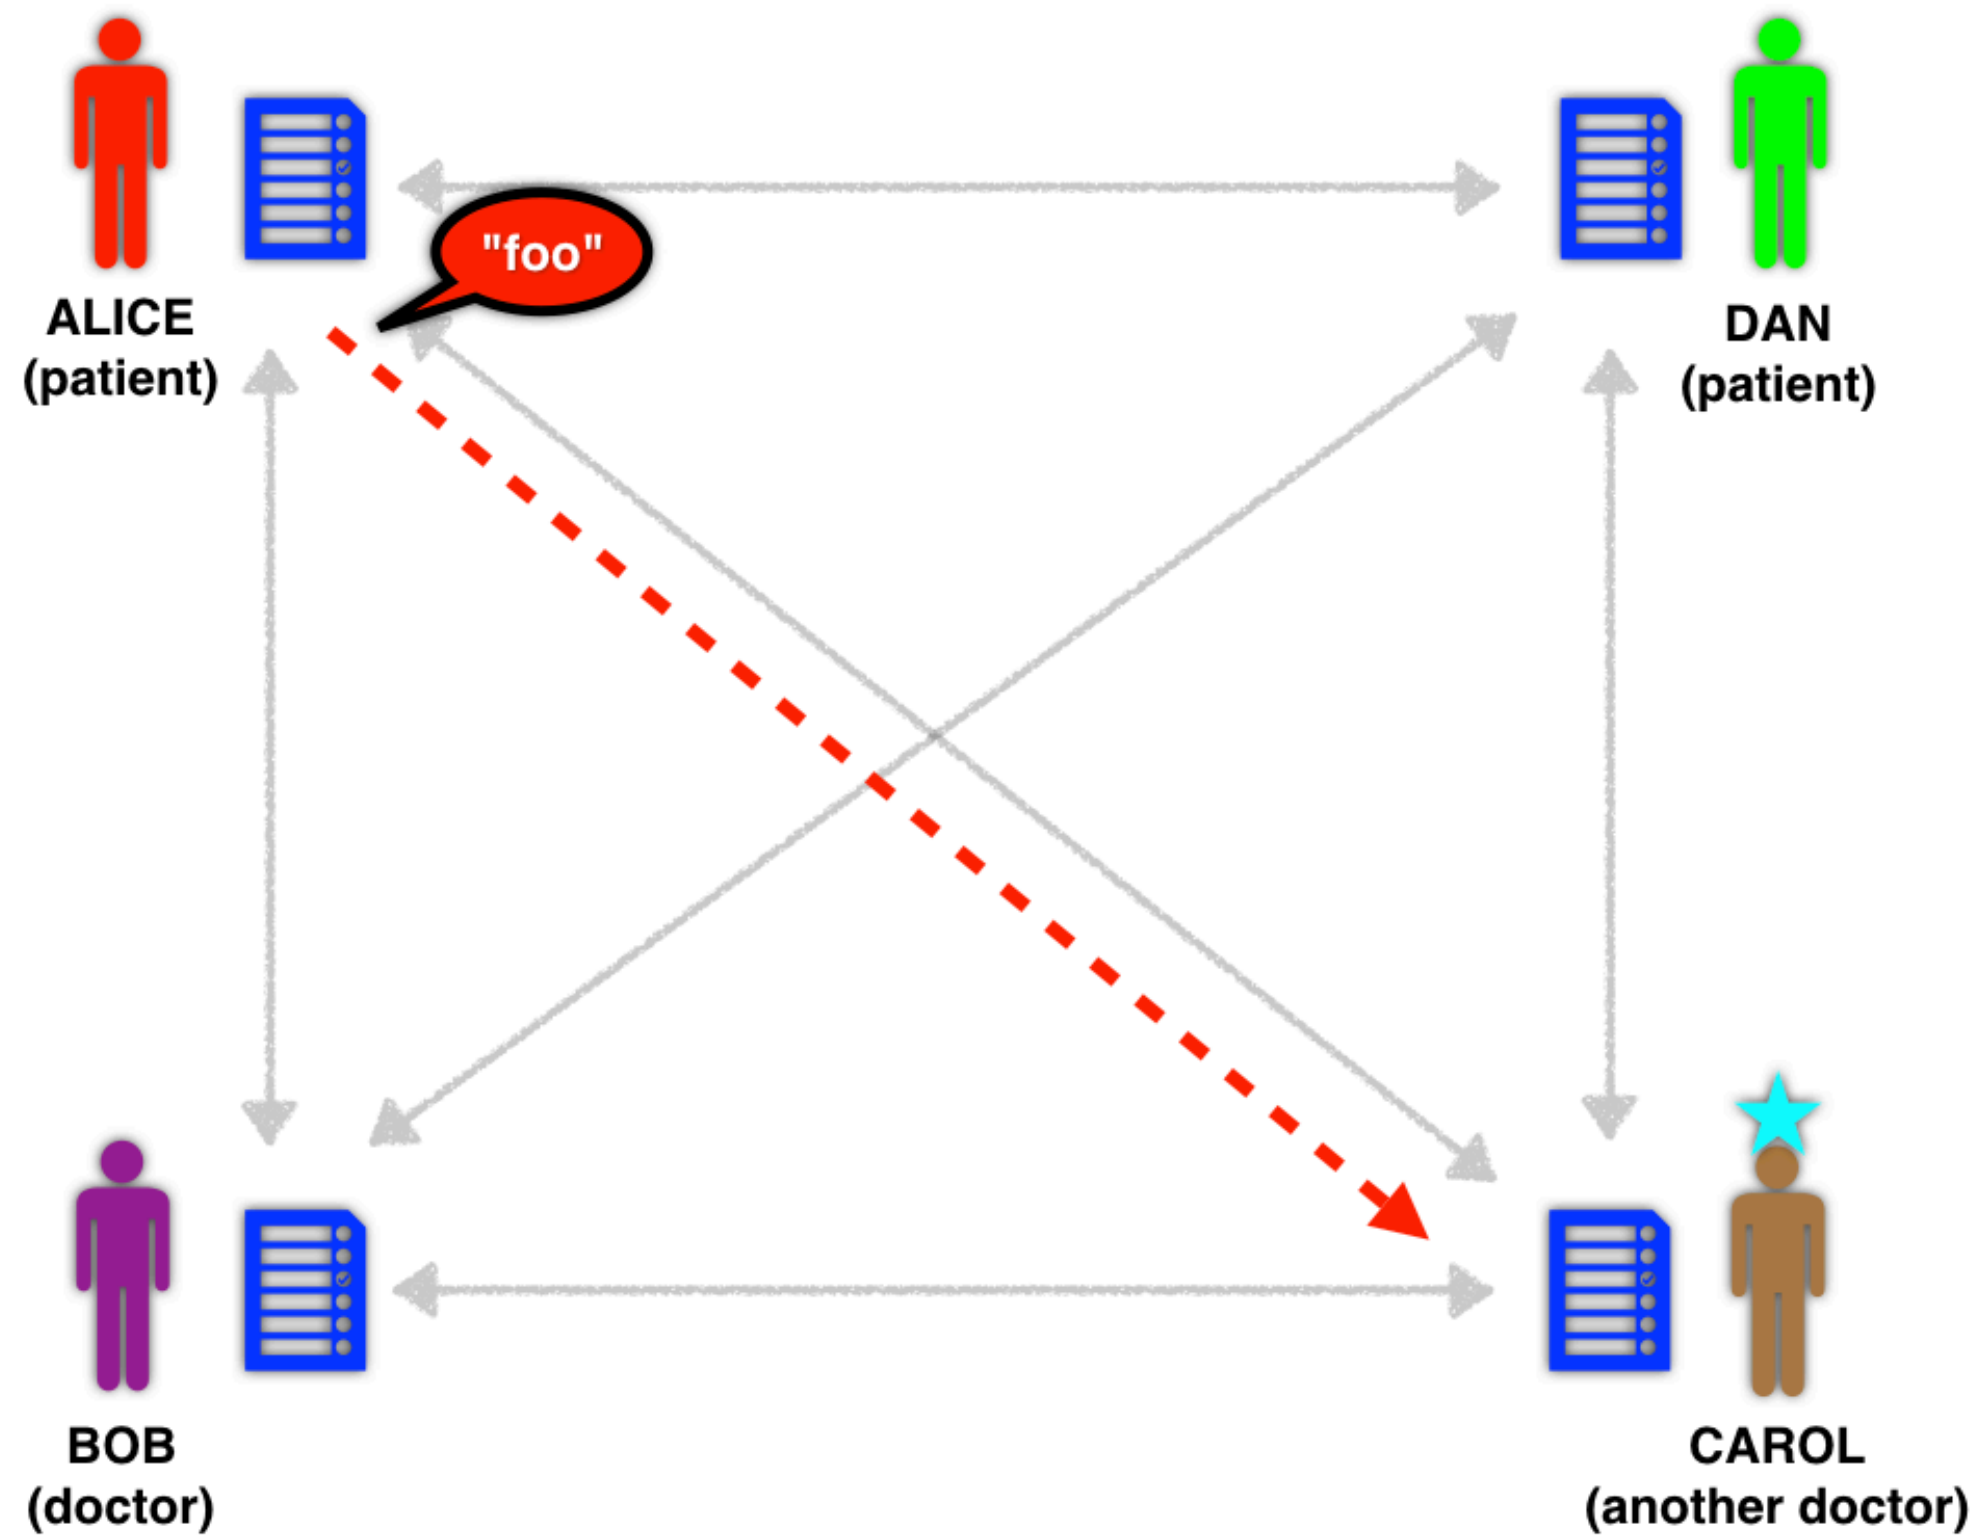
Let's start with collisions

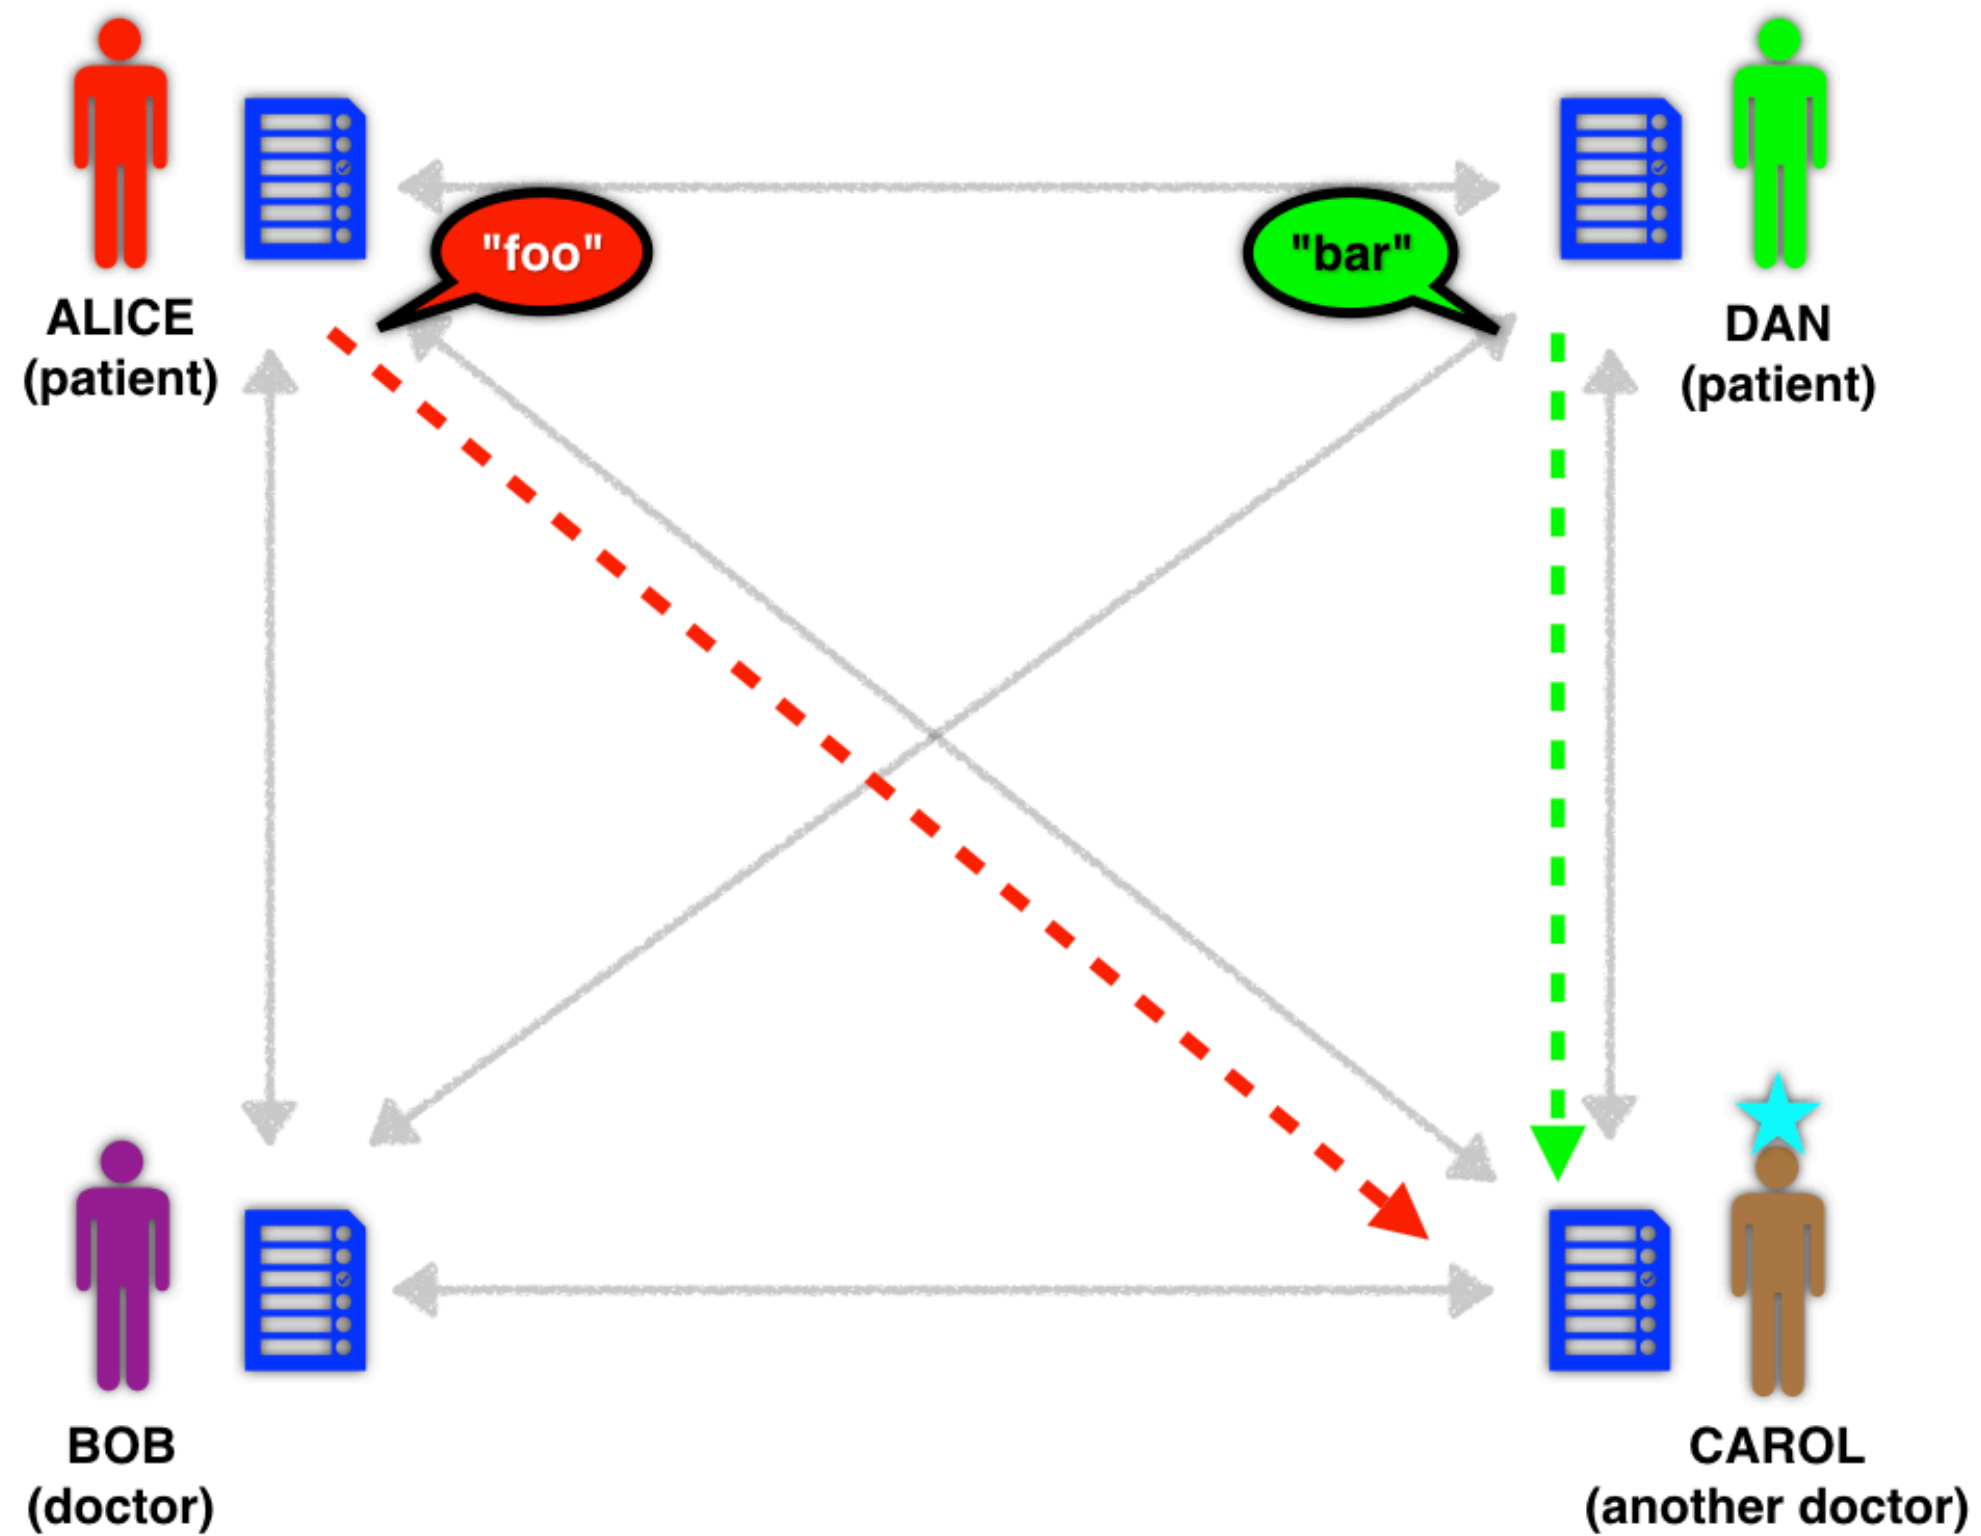
Log currently at: **[baz]**

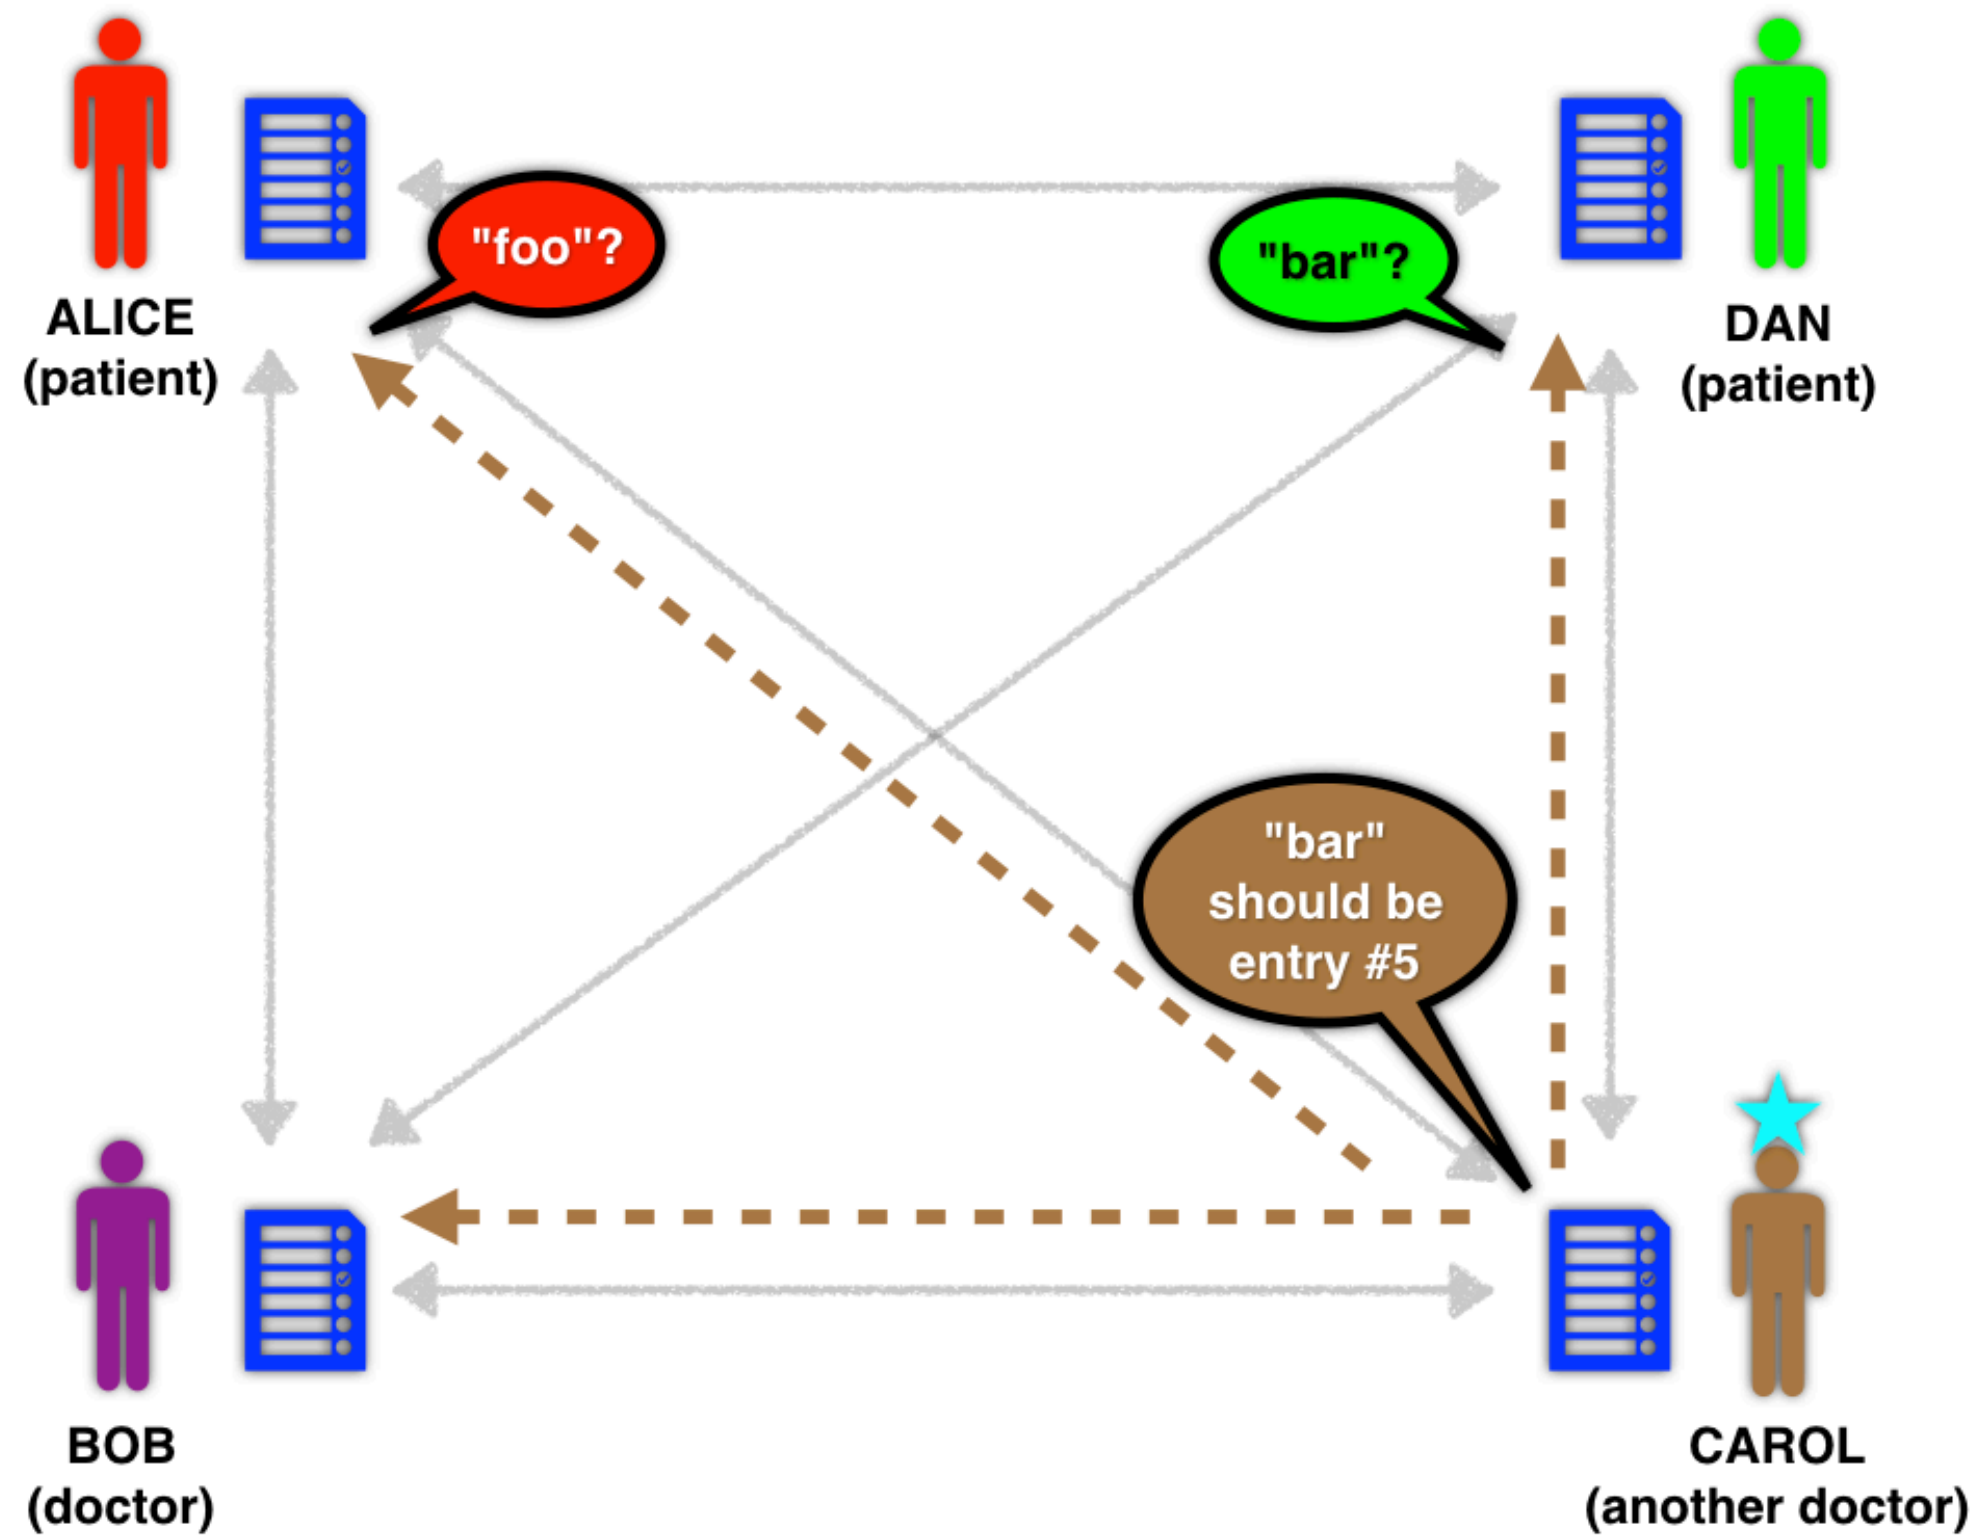
- Alice hashes baz, wants to add **foo**
- Bob hashes baz, wants to add **bar**

Who wins?

Ordering is complicated







Proper ordering flow

- Network assigns a leader who takes care of the ordering
- Alice and Bob no longer need to hash the log tip
- They just send their TXs (foo, bar) to the leader
- Leader decides on order arbitrarily (usually FIFO)
- Hashes log tip, broadcasts `<hash, orig_msg>` as usual
- Everybody goes along with it

In fact...

- If we just order one TX at a time, throughput will be low
- Leader waits to collect a batch of TXs and **orders the batch**
- So we now have a hashed chain of batches/**blocks**
- A **blockchain!** (Finally!)

A couple of more nodes on ordering

- Leader cannot fake transactions
- They still need to carry valid signatures
- If leader doesn't play nicely (byzantine), can be voted off
- Leader may change periodically

*Huge body of literature devoted to **distributed consensus***

Status check

1. Members have a copy of the transaction log
2. Members send signed writes to leader
3. Leader batches them, hashes previous block to establish order, then broadcasts this new batch to the network
4. Members inspect proposed batch, if everything looks good they adopt it as new addition to log

Privacy concerns

- Shared log holding all patient records: a privacy nightmare
- Don't store the records or reveal any patient identifying info
- Replace everything with hashes – use these as identifiers

Privacy concerns

- Shared log holding all patient records: a privacy nightmare
- Don't store the records or any patient identifying info
- Replace everything with hashes – use these as identifiers

For example

- Alice identified by the hash of her public key
- Alice's blood work represented simply by its hash

Wrong way to do it

Patient Name	Public Key	Record	Digital Signature	Who Can Read	Who Can Write
Alice	b83da6c0 9893ad	Blood work @ Duke Hospital (2016/05/ 01)	835cda89 7d74c3	Carol	Nobody

Right way to do it

Public Key	Record	Digital Signature	Who Can Read	Who Can Write
b83da6c0 9893ad	149dca74 ef983a	835cda89 7d74c3	fe45208d 673a67	Nobody

But *where* do you store the patient records?

- Use third-party **storage providers** built for this
- These services will have access to the blockchain
- They only process **signed requests** to read/write records
- They only accept them if the blockchain says that the associated public key is granted access to read/write on the blockchain

So how would it all work?

Coinbase-like services for your EHRs

Expect companies to spring up that:

- *Both* act as **storage providers** for your EHRs, *and*
- Help you **interact** with the blockchain network
- Like Coinbase for buying/selling Bitcoin
- For most folks, their only interaction with the Bitcoin blockchain is via the Coinbase website or mobile app

Who's actually holding a copy of the transaction log?

All **major** stakeholders:

- hospitals
- insurance companies
- these third-party Coinbase-like companies


Unrealistic to expect individuals (as shown before) to run their own nodes

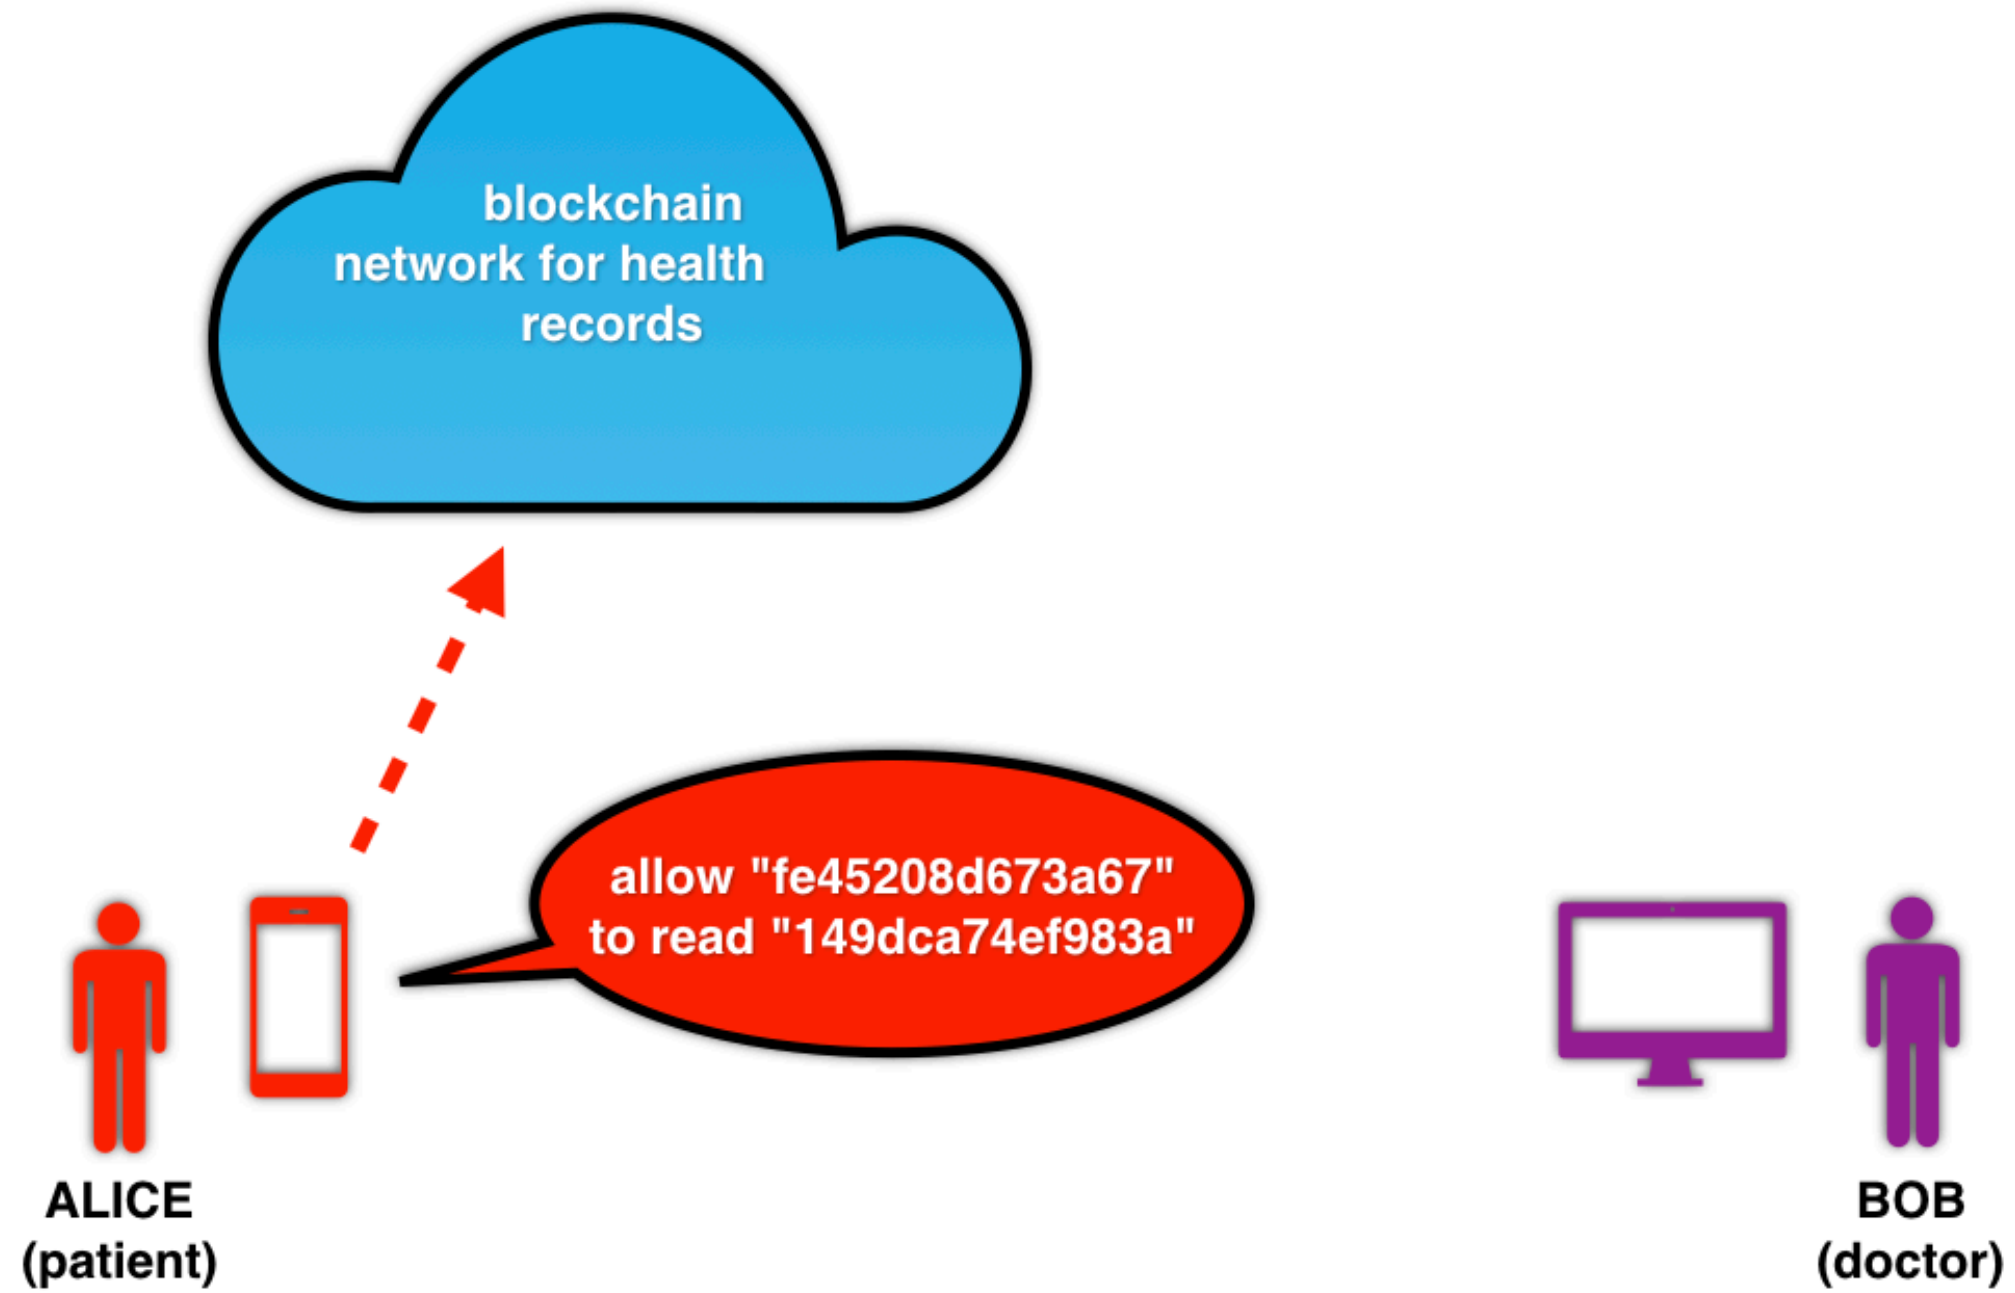
How does EHR data get into the system?


- Patients no longer get a hard copy of their lab work, but a digital one as well
- They upload these to their Coinbase-like service of choice
- They now have a record in that global table associated with their public key

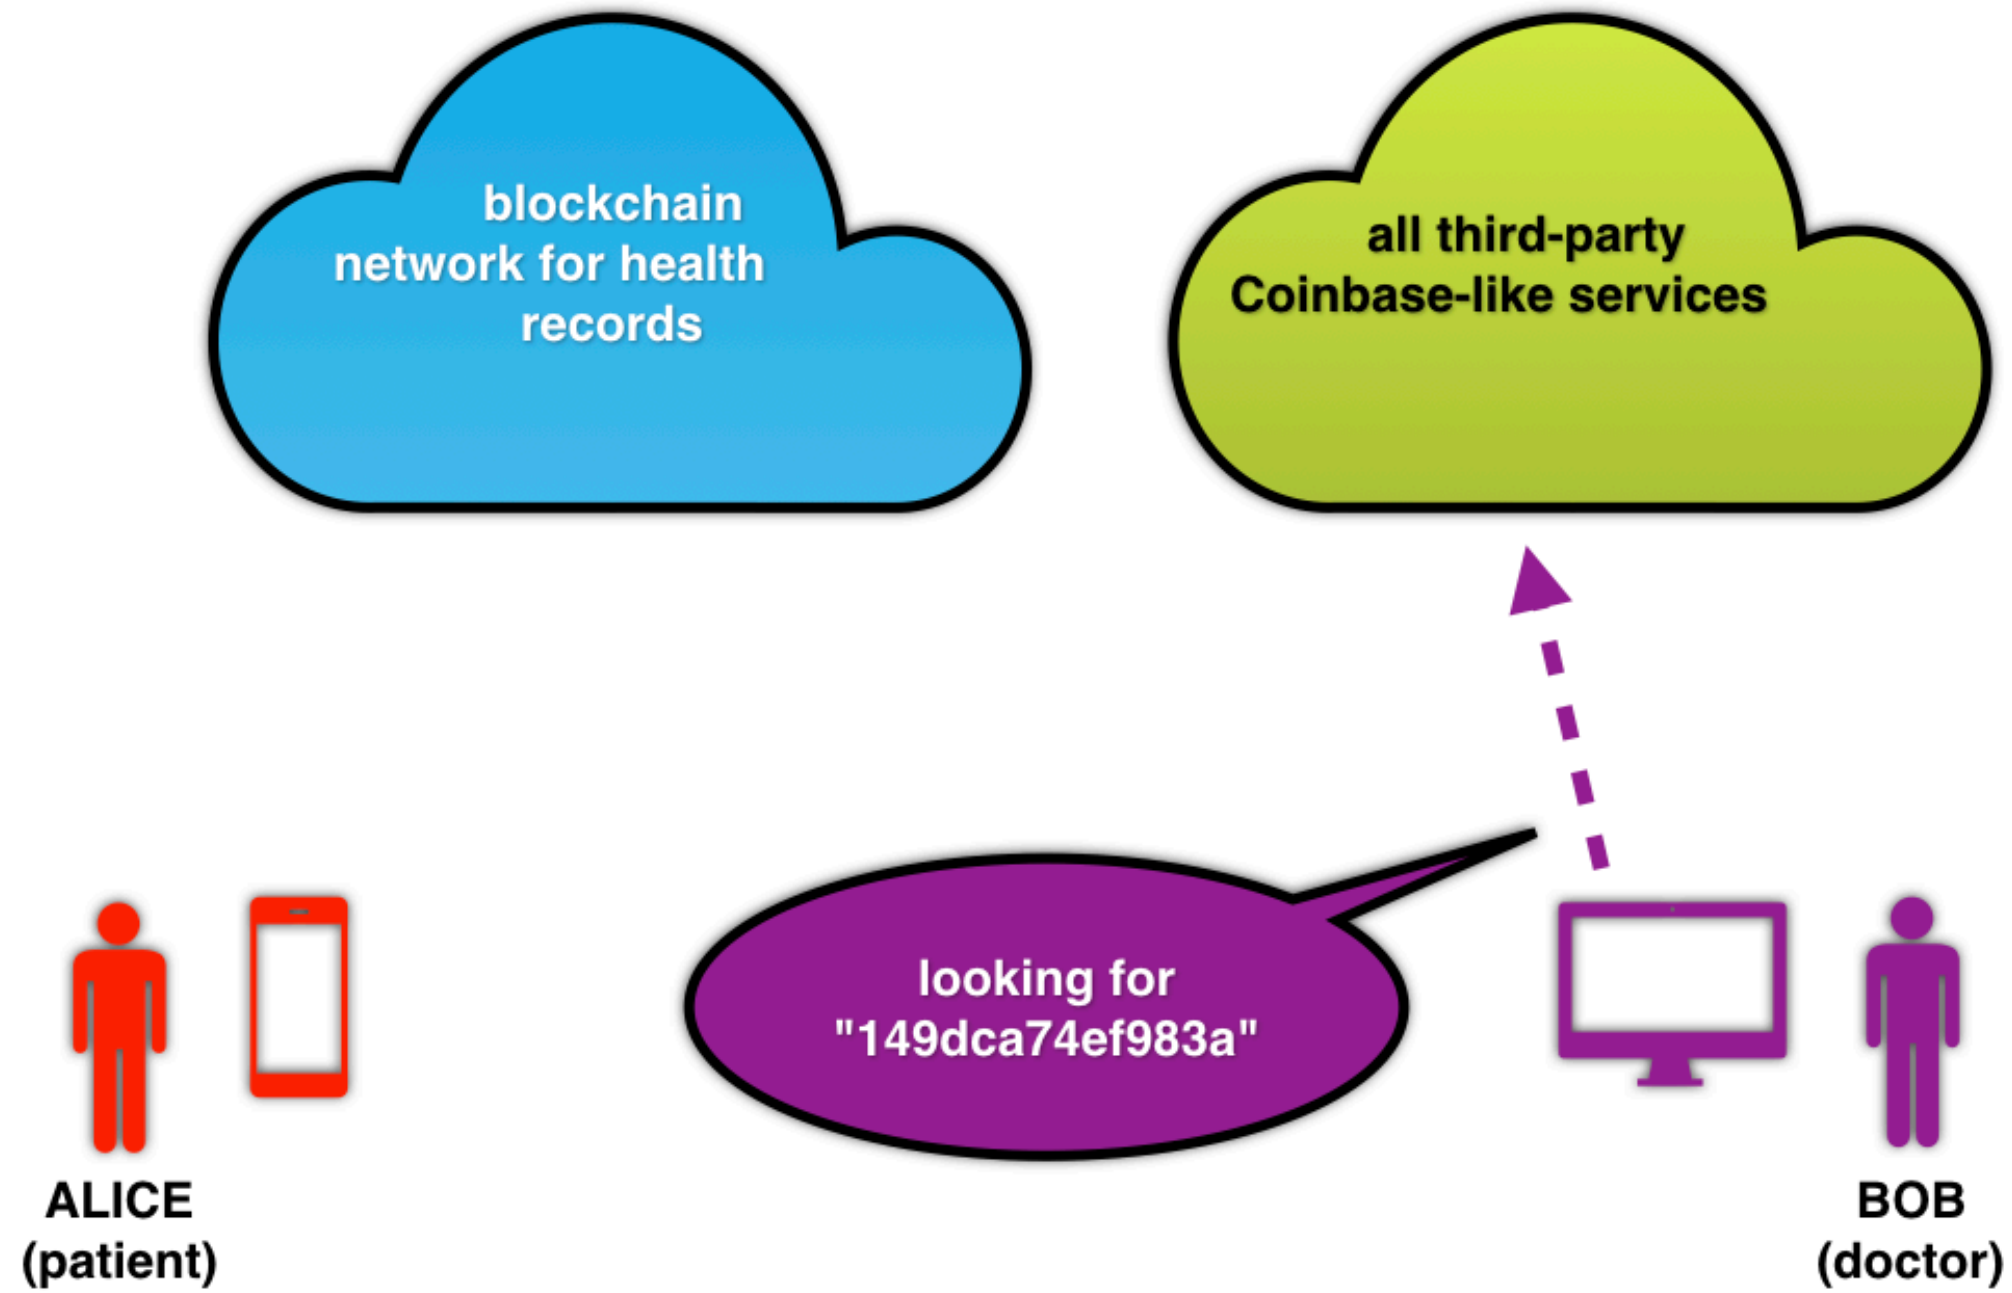
Expect heavy usage of mobile apps and QR codes here

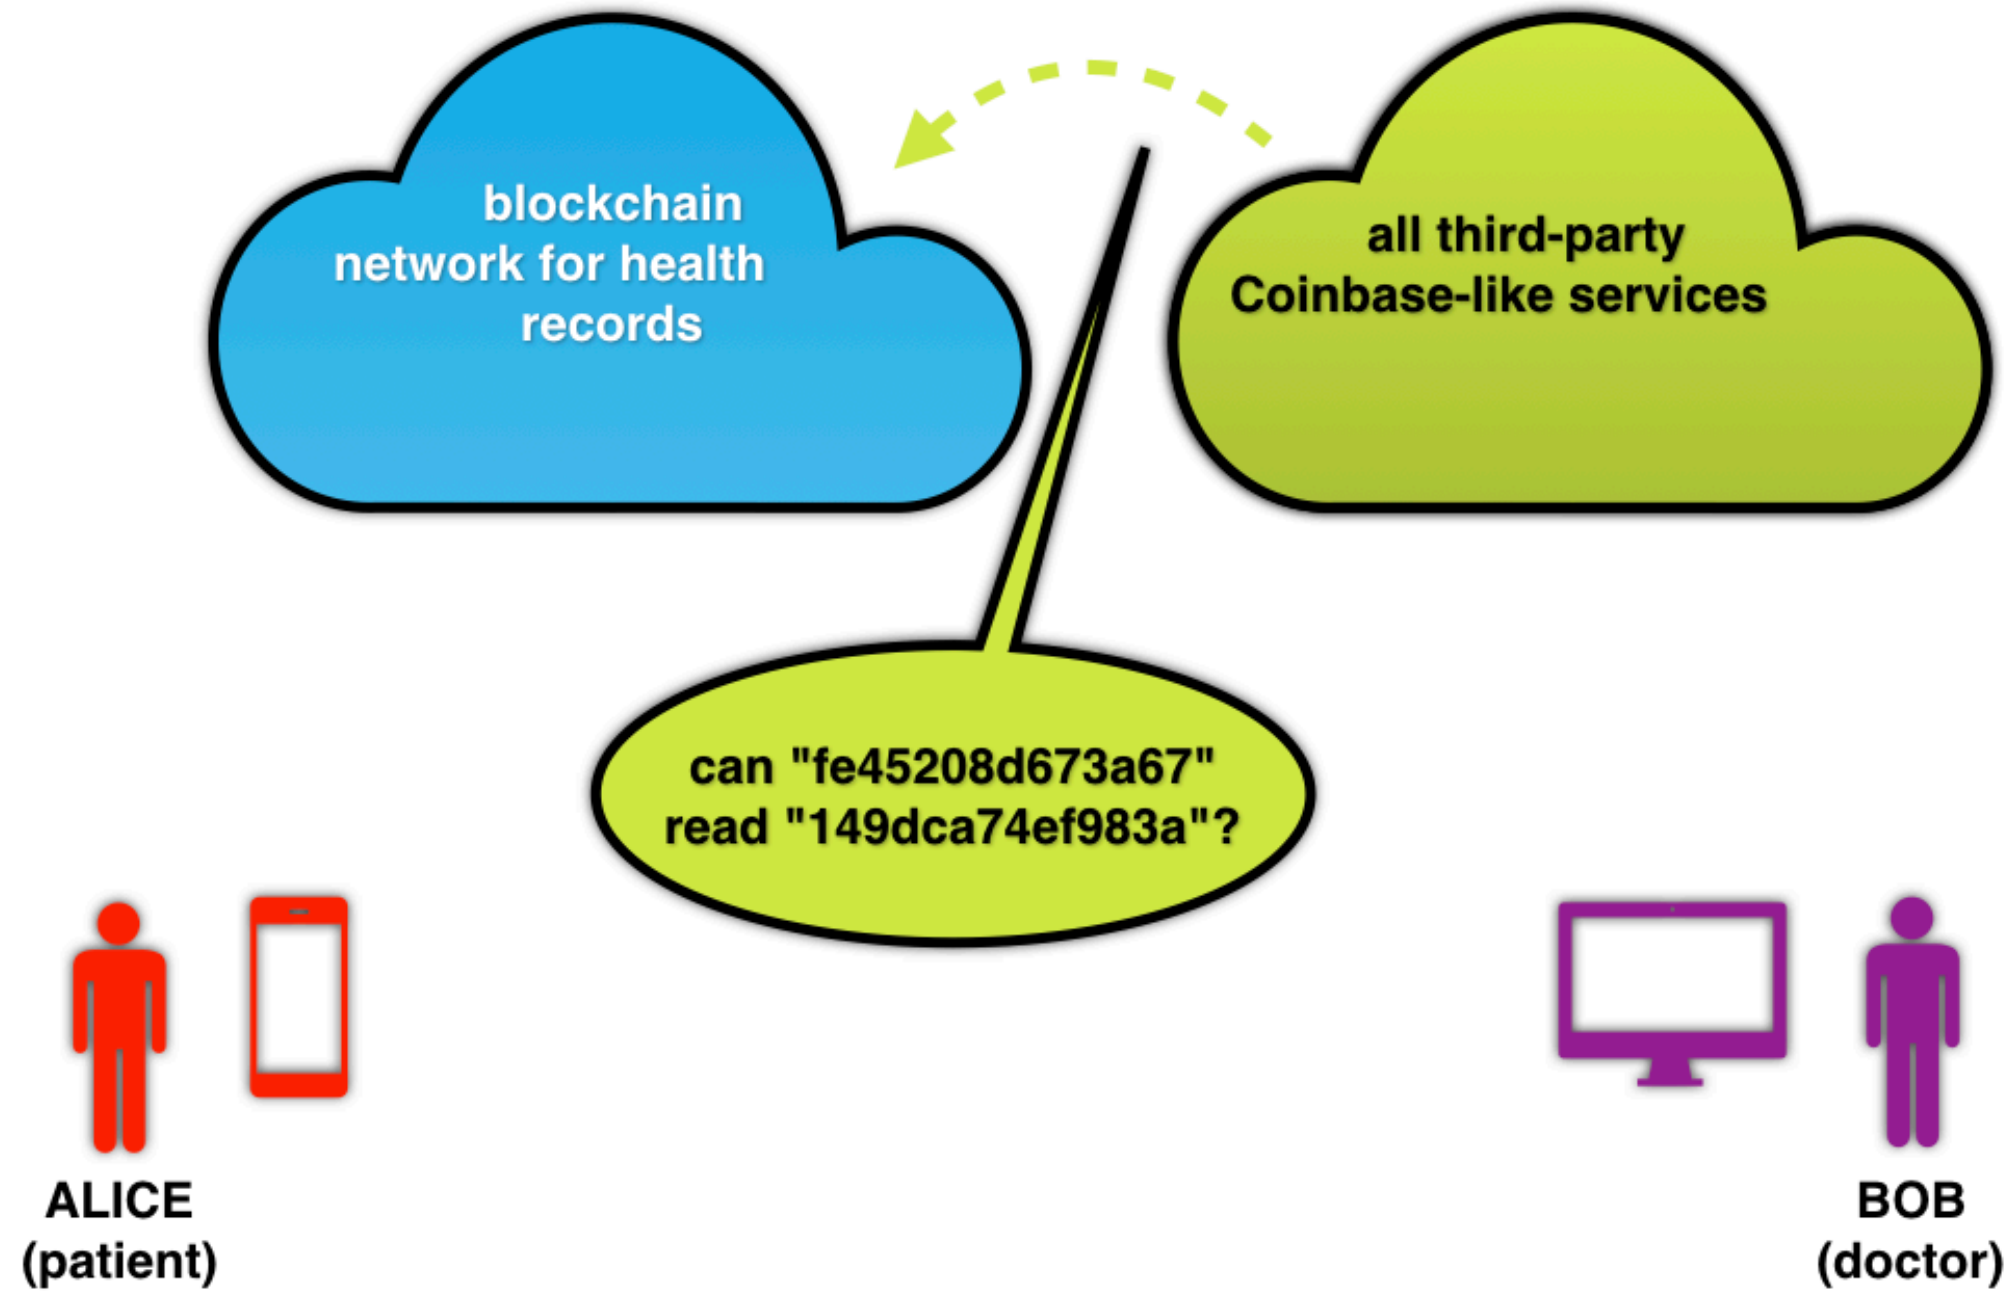
**Then a visit to the doctor looks
like this**

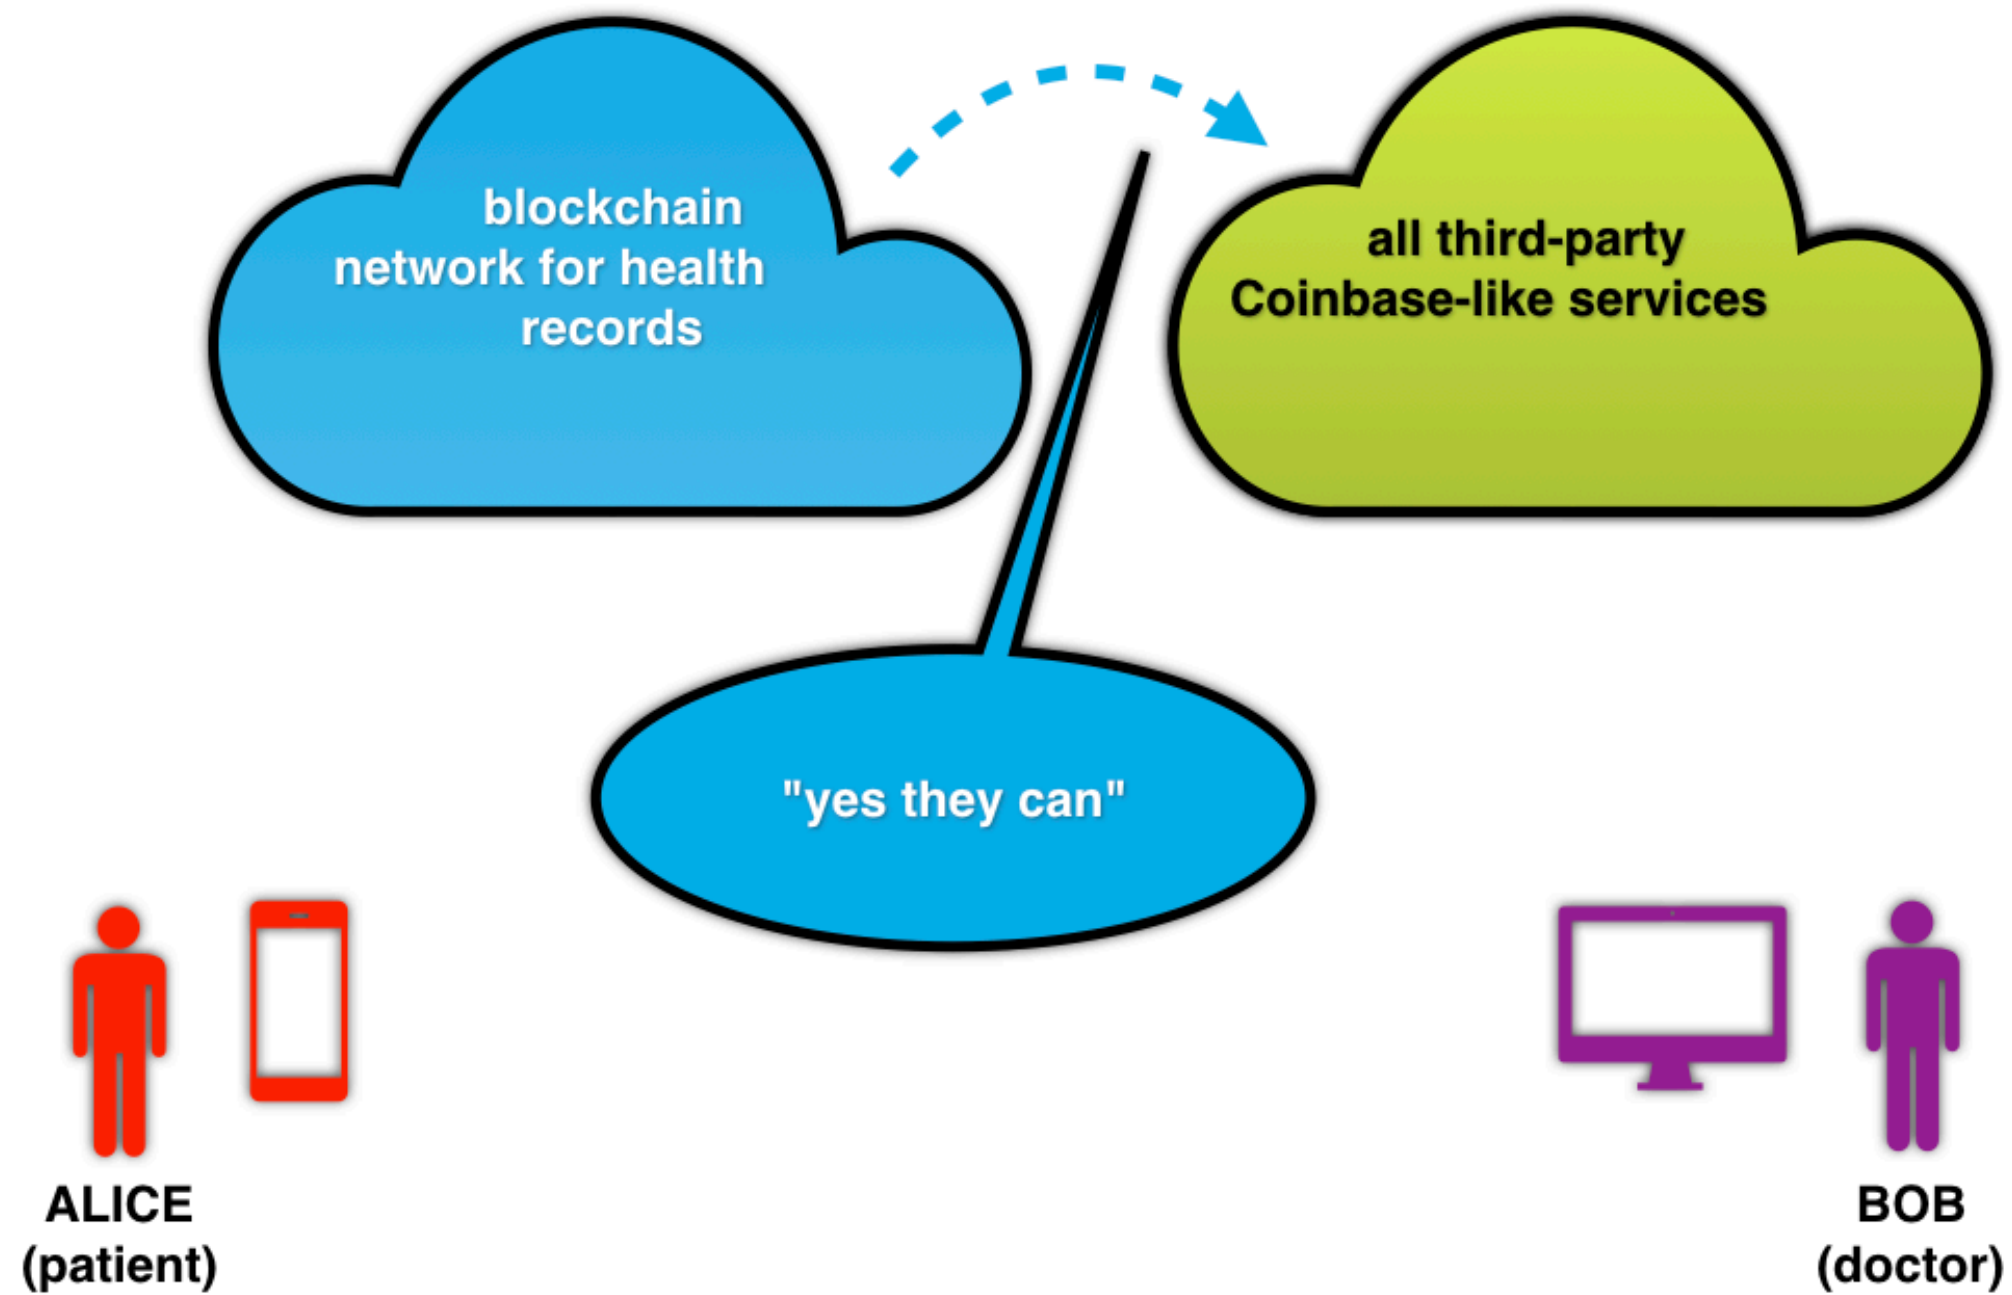


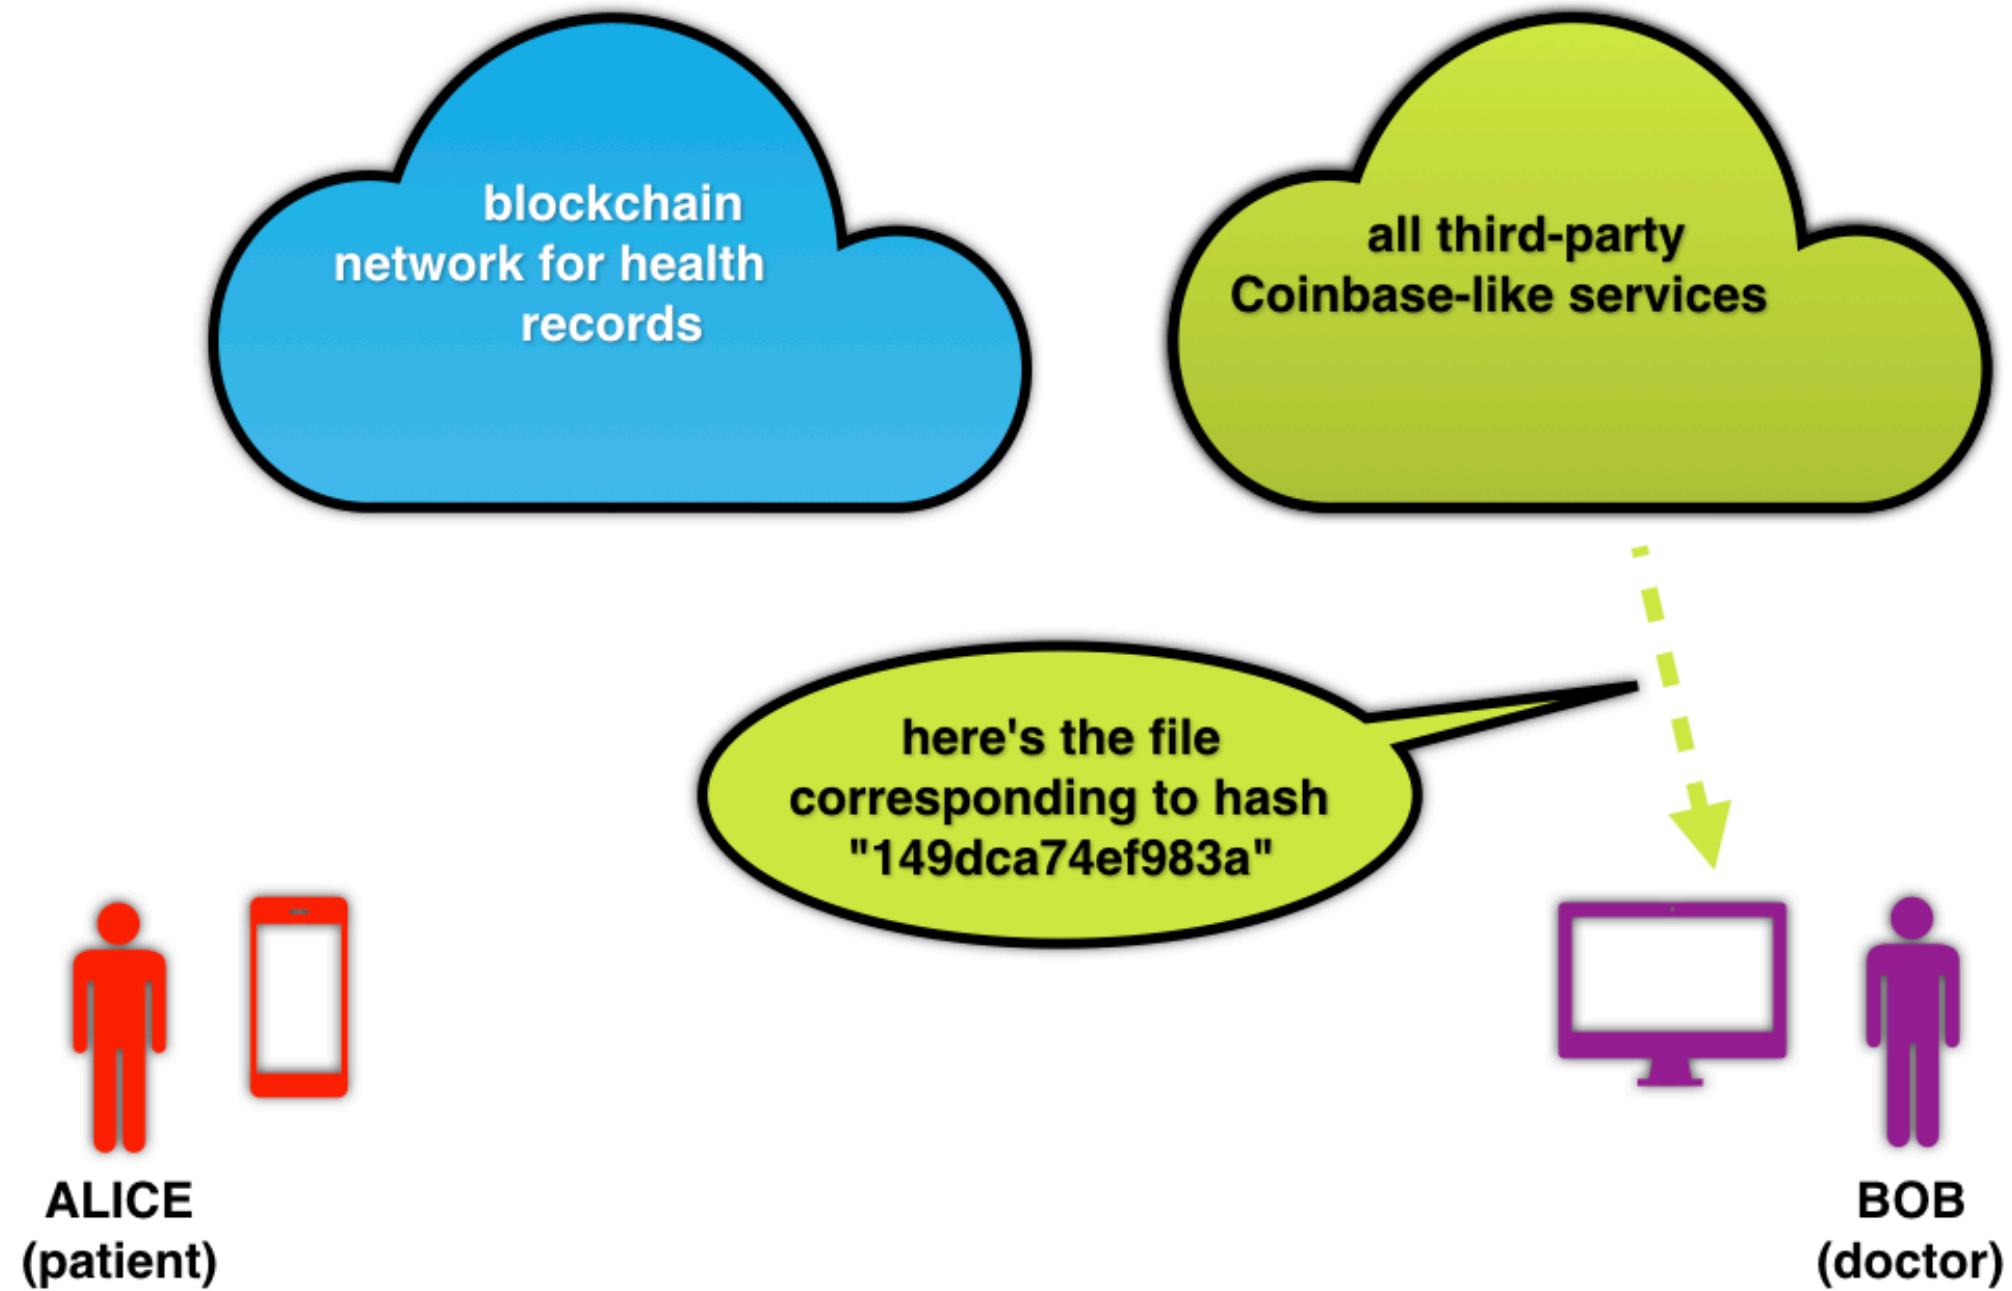












Visit to the doctor

- You visit the doctor & they need to see one of your records
- Using your mobile app, you:
 1. Scan the doctor's QR code (i.e. their public key)
 2. Then select the record you grant them access to
- Your mobile app then writes to the blockchain, updating the access list for the record

Visit to the doctor (cont.)

What the doctor's app does:

1. Reads latest blockchain updates
2. Observes TX adding them to the read-list for the record
3. Seeks that record (via its hash) in all participating third-party services

Visit to the doctor (cont.)

What the Coinbase-like third-party service then does:

1. Reads the blockchain to check if the request is allowed to access the record
2. If so, returns the file it stores locally to the reader

Epilogue

Blockchain

A replicated ledger of signed blocks forming a hash chain. (Also, a set of rules on how these blocks are added to the chain to begin with.)

- **Only way** for multiple non-trusting **writers** to share a DB
- Minimizes the need for trust
- Maximizes auditability, facilitates provenance tracking

Addendum

Smart contracts

- A misnomer: **simple scripts**
- Think of them as an "**if this, then that**" mechanism for your blockchain data
- Example: "if someone writes this data to my row on the DB, update this other row that I own"
- A way to create simple workflows *on the blockchain*

Questions?